

Argus: A General-Purpose Agentic Runtime for Long-Horizon Reasoning

Verification-Guided Persistence, Pivoting, and Runtime Self-Evolution

Argus Team



Shanghai Jiao Tong University

Abstract. Long-horizon reasoning is not repeated one-shot completion. Intermediate proofs, measured bottlenecks, failed routes, and specification contradictions can be valuable before a terminal artifact exists. Only when verified against explicit criteria can such signals justify refining the initial goal, constraint, or verifier.

We present Argus, a verification-centered runtime for persistent, pivot-capable long-horizon reasoning. Manager, Planner, Engineer, and Reviewer execute bounded missions over a durable task contract. Long-horizon reasoning is not merely persistence: verification determines when to continue, stop, replan, or surface evidence and alternatives for user clarification or authorization. Any resulting pivot is recorded in the durable contract, separating evidence-backed refinement from goal drift. Meanwhile, staged records of attempted approaches, accepted and rejected results, skills, verifiers, and routing update the runtime while model weights remain fixed. This *verifier-gated runtime self-evolution* improves how later missions navigate between the current goal and constraints. Between user decision boundaries, the harness remains fully autonomous.

In fully autonomous runs with no human intervention during execution, the same harness remains effective across seven GPT-5.5 benchmark arenas. It reaches approximately 78% (behind only Mythos & Fable 5) on SWE-Bench Pro versus 59% for Direct Copilot. After verifier-gated runtime evolution, mature SWE-Bench waves use 21% fewer solve input Tokens per task than startup waves. Across the broader suite, it reaches 76.8% on AARRI-Bench research tasks and a 28.0 gap on mathematical data synthesis, alongside competitive GPU-kernel and language-model-training results (Figure 1). Its artifacts extend beyond benchmark endpoints: an optimized RWKV6 kernel was merged into the upstream Flash Linear Attention repository, while the separately documented ACE-1 neural processing unit (NPU) remains operator-reported pending independent reproduction. A multi-day mathematical campaign retains one falsified route and six proof-backed frontier updates; six autonomous paper pipelines span 254 missions and 16 Stage rollbacks before reaching submission completion. These results show a general harness that can revise, recover, and accumulate verified approaches without forcing every long-horizon trajectory into a success claim.

Keywords: long-horizon reasoning; verified pivoting; user clarification; verification-centered agents; runtime self-evolution.

Project: argusbot.cn




Figure 1. Argus runtime and breadth evidence. The central diagram shows Manager authority over tasks, research verticals, and Stage transitions; Planner, Engineer, and Reviewer operate over a shared workspace whose knowledge, event log, artifacts, backlog, budget, daemon, and memory persist across bounded missions. The surrounding cards report outcomes from seven task-native evaluations. Their bars use independent scales and indicate comparison direction with arrows; they are breadth evidence rather than a single normalized leaderboard.

1 Introduction

Long-horizon reasoning is not a longer version of one-shot question answering. At the start of a mathematical investigation, neither the useful intermediate claims nor the reachable theorem frontier may be known. An engineering user may express a solution before understanding the actual performance, reliability, or integration constraint. In system verification, a failed proof obligation may reveal a defect in the implementation, an inconsistency in the specification, or a mistaken assumption shared by both. In each case, execution changes what the user knows about the problem itself.

This creates a different systems objective from autonomous completion of a fixed prompt. The system must preserve a standing intent while allowing the operational goal, constraints, and acceptance criteria to become more precise. It must value verified intermediate progress—a counterexample, a stronger bound, a measured bottleneck, a minimal failing trace, or a corrected specification—even when no terminal artifact yet exists. It must also expose these updates to the user, whose judgment can revise priorities and constraints for the next mission. The system challenge is therefore not merely to persist, but to determine when evidence supports continuing the current route and when it warrants a pivot in approach, objective, constraints, or verification criteria.

Large language model agents have progressed from tool use and software execution to complete research loops. ReAct, SWE-agent, and OpenHands established practical interfaces between models and external environments (Yao et al., 2023; Yang et al., 2024; Wang et al., 2024a). The AI Scientist, CycleResearcher, AutoSci, FARS, and Arbor extend this interaction into iterative research programs covering different combinations of planning, experiments, review, memory, and scientific communication (Lu et al., 2024; Weng et al., 2025; Qian et al., 2026; Tang et al., 2026; Jin et al., 2026). Their progress makes the central systems problem clear: long execution becomes reasoning capability only when verified results accumulate across iterations and the runtime can change direction when those results invalidate the current route or task formulation.

In practice, long-running agents fail in four recurrent ways, consistent with results showing that agent success declines as software-task horizons grow (Kwa et al., 2025). First, *continuity* is fragile: a growing transcript is eventually compacted or dropped, and later calls spend time reconstructing decisions that were already made. Second, *acceptance* is ambiguous when the component that performs the work also declares it complete. Third, *experience* is lost when only the final artifact survives; failed branches, contaminated measurements, and useful procedures disappear with the session that produced them. Hierarchical and workflow-oriented memory systems address parts of this problem (Packer et al., 2023; Wang et al., 2024b; Xu et al., 2025). A larger context window extends session continuity; durable institutional memory additionally requires explicit state, ownership, and update rules. Fourth, the task contract itself can be wrong or incomplete. A runtime that treats the initial prompt, specification, or evaluator as immutable may optimize the wrong quantity with increasing efficiency. Conversely, a runtime that freely rewrites its objective turns learning into unauditable goal drift. The missing mechanism is verified contract refinement: evidence must identify why a goal or constraint is insufficient, and an authorized decision must record what changed and what remains invariant.

We present Argus (Autonomous Research Generation and Understanding System), a verification-centered runtime for persistent and pivot-capable long-horizon reasoning. Its central abstraction is a sequence of *bounded missions* executed against a durable, revisable task contract. Four model-driven roles divide authority: a Manager anchors the objective and campaign state, a Planner selects the next units of work, an Engineer implements and evaluates them, and a Reviewer inspects artifacts and run outputs before issuing the completion verdict or identifying a contract-level contradiction. The user remains outside the repetitive round loop but retains authority at consequential boundaries. Manager-visible evidence, alternatives, and unresolved questions make clarification possible when the correct goal or constraint cannot be settled by artifacts alone. These roles operate across three planes—control, execution, and records—so that scheduling, work, and record keeping remain separable even though they participate in one continuous loop.

This report describes the system and evaluations available through August 4, 2026; historical benchmark traces are reported under the configurations that produced them.

The runtime is designed to accumulate two coupled forms of state without updating the base model. First, it self-evolves the *approach* between the current goal and constraints: each Stage retains attempted routes, measurements, Reviewer verdicts, accepted and rejected results, reusable skills, verifiers, and routing decisions, so later missions begin from a changed search policy rather than merely a longer transcript. Second, verified evidence can sharpen the user-facing task contract: goals can be decomposed or reframed, constraints can be discovered or corrected, and acceptance tests can be strengthened. Verifier-gated runtime self-evolution is the engine that improves execution under the current contract. Verification-guided pivoting changes that contract only when evidence exposes a material mismatch; user clarification or authorization governs changes that cannot be settled by task-native evidence alone. The term *dense intelligence* describes the corresponding design goal: keep decision, execution, and verification connected over the research horizon rather than concentrating useful reasoning into isolated bursts. Section 3 formalizes this conversion and the experiments measure its operating factors.

We sharpen this design language into a verified-progress theory. The final artifact is a projection of a richer typed trajectory; intermediate states can carry task-native value; reviewed compression retains the parts needed for later decisions; and an accepted update counts as compounding only when it reduces future task risk or resource cost. Contract refinement is admissible only when it preserves the standing intent, cites new evidence, records the changed assumptions, and passes the appropriate user

or Manager authority boundary. This separates a justified pivot from goal drift.

We evaluate the resulting system across seven benchmark arenas covering software repair, GPU kernels, language-model training, training speed, research-assistant tasks, and mathematical data synthesis. Beyond the headline SWE-Bench result, the suite includes 76.8% on AARRI-Bench and a 28.0 mathematical-data gap, both above their reported references. SWE-Bench Pro is one row in this benchmark suite; its 731-task trajectory additionally enables analyses of Skill/Wiki evolution and Reviewer intervention. A de-duplicated inventory records 41 research artifacts across six programs as a separate breadth measure. We additionally report one mathematical campaign as a representative vertical trace, showing how the four roles hand off the same objective through recovery, route testing, proof production, revision, acceptance, and retained state; it preserves one falsified route and six proof-backed theorem-frontier updates. Finally, we reconstruct six Argus paper-production campaigns spanning 254 missions and 16 Stage rollbacks from their structured Stage and role logs. These traces are especially relevant to pivoting because several initial positive hypotheses fail and become narrower diagnostic or negative-results contributions through evidence-backed rollback.

This report makes seven contributions:

- C1.** We formulate long-horizon reasoning as verified progress over a revisable task contract $K_t = (i, o_t, c_t, v_t)$ comprising standing intent, objectives, constraints, and verification criteria, rather than completion against an immutable prompt.
- C2.** We introduce a role-separated runtime for verification-guided persistence and pivoting that combines durable intent, bounded execution, independent review, user-visible contract refinement, append-only records, and fixed-model runtime evolution.
- C3.** We report one unified table across seven benchmark arenas, preserving each benchmark’s backbone, backend, protocol, native metric, and reference.
- C4.** We analyze the SWE-Bench Pro row in depth: Argus reaches approximately 78% accuracy versus 59% for Direct Copilot at $1.41\times$ aggregate Tokens; mature Waves are more efficient, and Reviewer intervention yields 34 verifier recoveries and 22 strict review-loop rescues.
- C5.** We reconstruct one mathematical campaign as a role-resolved Agent trace, showing why falsified routes and intermediate theorem-frontier updates are valuable before a terminal proof and how they reshape subsequent work.
- C6.** We present a six-project autonomous paper-production case study covering 640 campaign-hours, 576 Engineer rounds, 286 Reviewer revisions, and 16 Stage rollbacks, with all six canonical pipelines reaching submission completion. A representative campaign uses seven verified no-go decisions to reframe a failed method search as a 4,500-row negative-results study, then repairs two late submission defects without discarding the accumulated research state.
- C7.** We document two artifact-level outcomes: a human-reviewed TileLang kernel merged into Flash Linear Attention and the operator-reported ACE-1 NPU prototype, while keeping their distinct evidence levels explicit.

The remainder of the report follows a conventional research-paper organization. Section 2 positions the system against prior agent and autonomous research work. Section 3 formalizes the operating problem; Section 4 presents the runtime; Sections 5 and 6 describe the evaluation protocol and results; and Sections 7–9 discuss interpretation, limitations, and conclusions. The appendix contains detailed experimental tables and notation.

2 Related Work

2.1 Autonomous Research Systems

Autonomous research systems range from tightly scoped experiment loops to complete research lifecycles. Karpathy’s *autoresearch* repeatedly modifies, trains, and evaluates a model under a fixed compute budget (Karpathy, 2026); AIDE organizes iterative code search for machine-learning improvements (Jiang et al., 2025). FunSearch and AlphaEvolve extend evaluator-guided program search to mathematical, algorithmic, and systems problems (Romera-Paredes et al., 2024; Novikov et al., 2025). These systems show how a clear objective and an executable evaluator can turn model calls into sustained search.

Full-lifecycle systems add problem selection, planning, experimentation, review, and scientific communication. The AI Scientist introduced an end-to-end research pipeline, while AI Scientist-v2 added agentic tree search and produced a workshop-accepted paper (Lu et al., 2024; Yamada et al., 2025). Agent Laboratory organizes specialized roles across the research process (Schmidgall et al., 2025), and CycleResearcher closes the loop between automated research and automated review (Weng et al., 2025). More recently, AutoSci proposed a memory-centric system spanning the complete scientific lifecycle, and FARS reported large-scale deployment of fully automated research workflows across 67 fine-grained AI/ML topics (Qian et al., 2026; Tang et al., 2026).

A complementary line focuses on scientific ideation and collaboration. ResearchAgent iteratively generates and reviews ideas grounded in the literature (Baek et al., 2024); SciAgents combines knowledge graphs with multi-agent reasoning for materials discovery (Ghafarollahi and Buehler, 2024); and Co-Scientist uses a multi-agent generate–debate–evolve loop for biomedical hypothesis discovery (Gottweis et al., 2025). AgentRxiv studies cumulative research through a shared repository of prior work (Schmidgall and Moor, 2025), while Arbor uses a persistent hypothesis tree and isolated executors to refine long-running research programs (Jin et al., 2026).

These systems are strongest when the optimization target and evaluator are already trusted. Long-horizon expert reasoning often begins one step earlier: deciding which target, constraint set, and verifier correctly express the real problem. Argus uses many mechanisms developed by this autonomous-research family but targets a different center of gravity. Its primary object is not an autonomous run against a fixed objective; it is the verified relationship among a user, a revisable task contract, and an evolving runtime. The framework carries accepted results, failed routes, tools, skills, task definitions, and contract revisions across bounded missions.

2.2 Human–Agent Collaboration and Task Refinement

Many agent systems model the human as the source of an initial instruction, an occasional approver, or the recipient of a final artifact. That interface is insufficient when task discovery is part of the work. Users contribute tacit preferences, domain constraints, risk judgments, and authority over acceptable tradeoffs. Agents contribute rapid search, implementation, measurement, and a larger explicit process record. Neither side necessarily possesses the complete task specification at the beginning.

Our focus is verification-guided pivoting: verified agent work can expose a false assumption or missing constraint, while user clarification or authorization can change the goals, constraints, verifiers, or priorities governing later agent work. This is not unrestricted conversational adaptation. A contract update must be traceable to evidence, preserve declared invariants, expose alternatives and uncertainty, and pass an authorized decision boundary. Argus places verification at the center because verification is the shared language through which both parties can distinguish discovery from drift.

This positioning differs from three nearby patterns. A general agent harness mainly extends tool use and execution reliability. A fixed-evaluator autoresearch loop mainly improves a candidate under a known score. A fully autonomous research pipeline mainly reduces human intervention across a predefined lifecycle. Argus can support all three, but its distinctive objective is to help users discover and reach better-specified goals through persistent, review-gated work.

2.3 Long-Horizon Agent Runtimes and Memory

Tool-using agents established the basic interaction loop. Toolformer learns when to invoke external tools (Schick et al., 2023), ReAct interleaves reasoning with environment actions (Yao et al., 2023), and SWE-agent and OpenHands expose repository-scale execution through agent–computer interfaces (Yang et al., 2024; Wang et al., 2024a). These systems provide the execution substrate for long tasks; autonomous research additionally requires continuity across many sessions, experiments, and changes of direction.

Persistent-memory systems address this continuity from several angles. Generative Agents combine stored observations with reflection and retrieval (Park et al., 2023); Reflexion retains feedback across attempts (Shinn et al., 2023); Voyager accumulates executable skills (Wang et al., 2023); and MemGPT manages multiple memory tiers (Packer et al., 2023). Agent Workflow Memory induces reusable procedures from past trajectories (Wang et al., 2024b), while A-MEM organizes memories through agentic linking and retrieval (Xu et al., 2025). Argus combines these ideas with explicit ownership: execution produces candidate updates, review commits the retained form, and the complete event stream remains separate from the bounded working checkpoint.

2.4 Learning from Research Trajectories

Intermediate trajectories are also valuable training material. Process supervision can train stronger mathematical verifiers than outcome-only supervision (Lightman et al., 2023). AgentInstruct generates post-training corpora through agentic flows, Agent-FLAN studies data and tuning strategies for agent behavior, and Agent Lightning converts existing agent trajectories into reinforcement-learning transitions (Mitra et al., 2024; Chen et al., 2024; Luo et al., 2025). Test-time-compute studies further show that search and verification should be allocated according to task difficulty rather than by a fixed sampling rule (Snell et al., 2024).

The trajectories retained by Argus connect model outputs to concrete states, tool actions, measurements, artifacts, review decisions, and runtime updates. They therefore serve two purposes: coordinating the next mission online and forming a structured corpus for later supervised, preference-based, or reinforcement learning.

2.5 Evaluation of Research Agents

Software and general-agent benchmarks established repository-level issue repair and multi-environment analytical evaluation as core Agent capabilities (Jimenez et al., 2024; Liu et al., 2023; Ma et al., 2024). Research-agent benchmarks increasingly test complete research behavior rather than isolated question answering. MLE-bench packages Kaggle competitions as machine-learning engineering tasks (Chan et al., 2024); RE-Bench compares AI agents with human experts in open-ended research engineering (Wijk et al., 2024); AARRI-Bench evaluates granular tasks from the research lifecycle (Wang et al., 2026); and AIRS-Bench targets frontier research-science agents across multiple disciplines (Lupidi et al., 2026). Together they motivate evaluation in task-native units, with explicit resource accounting and trajectory-level analysis.

This report combines that benchmark view with a continuous system trace. The benchmark suite measures outcomes across

seven arenas, while the SWE-Bench Pro and mathematical studies expose runtime evolution, review-driven correction, and cross-mission research progress.

3 Problem Formulation

3.1 Long-Horizon Tasks as Evolving Contracts

The public Argus formulation defines a *dense-intelligence task* as one that sustains reasoning, tool use, verification, and iteration across a continuous time window until it produces measurable progress (Argus Team, 2026). Three conditions determine whether a task has this shape:

- T1. Invention.** The answer is not directly retrievable from a manual or database; it must be discovered through search and feedback.
- T2. Long horizon.** One pass is insufficient. Progress requires many dependent iterations over hours or days.
- T3. Progressive verifiability.** Task-native evidence can distinguish a genuine frontier update from activity or a persuasive but incorrect claim, even when terminal success is not yet available.

The defining loop is therefore proposal $\rightarrow$ execution $\rightarrow$ measurement $\rightarrow$ verification $\rightarrow$ contract refinement. The last transition is essential. Measurement may reveal that the proposed method is wrong, but it may also reveal that the stated objective is underspecified, a constraint is incompatible with the desired outcome, or the evaluator rewards the wrong behavior. Performance engineering, system verification, scientific search, and mathematical research differ in domain but share this operational structure.

We represent the working task at mission t by a contract

$$K_t = (t, o_t, c_t, v_t), \quad (1)$$

where t is the standing user intent, o_t is the current operational objective, c_t is the set of known constraints and assumptions, and v_t is the current verification procedure and acceptance criterion. The standing intent is durable; the other components are hypotheses about how to realize and recognize it. This distinction lets a campaign learn that a benchmark, specification, or requested implementation is wrong without silently abandoning why the user started the work.

A campaign also contains an initial artifact y_0 , a resource budget B , an user epistemic state X_t , and runtime state H_t . X_t denotes the user’s current understanding of feasible routes, tradeoffs, uncertainty, and desired constraints; it is not a psychological model inferred from private behavior, but the explicit decisions, clarifications, and priorities made visible at the interface. The artifact may be code, a model, a dataset, an experimental harness, a proof, a specification, or a manuscript. The campaign is divided into bounded missions indexed by t , each producing

$$\tau_t = (K_t, s_t, a_t, y_t, e_t, r_t), \quad (2)$$

where s_t is the mission state, a_t the actions taken, y_t the resulting artifact state, e_t the observed measurements, and r_t the review outcome. A mission must terminate, pause, or transfer control explicitly; continuity belongs to the persistent campaign state rather than an unbounded model session.

3.2 Verified Contract Refinement, Not Goal Drift

Evidence can propose a new contract K'_t , but change is admissible only through a guarded transition

$$K_{t+1} = \mathcal{A}(K_t, K'_t, e_t, r_t, u_t), \quad (3)$$

where u_t is the authorized expert or Manager decision. The admission operator $\mathcal{A}$ accepts a refinement only when four conditions hold:

- R1. Evidence.** The proposed change cites a measurement, proof obligation, counterexample, failing trace, artifact inspection, or explicit user clarification.
- R2. Intent preservation.** The update explains how the new objective, constraint, or verifier better serves t , or explicitly records an user-authorized change to that standing intent.
- R3. Authority.** Reviewer can reject unsupported claims and recommend replanning; Manager controls campaign and Stage transitions; the user controls value judgments and material changes that cannot be settled by task-native evidence alone.
- R4. Provenance.** The old and new contracts, triggering evidence, unresolved uncertainty, and decision owner remain inspectable.

Goal drift changes what is optimized without a justified authority and evidence chain. A justified pivot changes the operational contract because verified work has made the original formulation more precise, less contradictory, or more faithful to the standing intent. A rollback is therefore not automatically regression, and preserving the exact initial wording is not automatically alignment.

3.3 Intermediate Progress as Task-Native Value

Terminal success is only one source of value. Let F_t denote the accepted frontier—the set of claims, artifacts, constraints, and excluded routes that have survived the current verifier. We write mission value as

$$R_t = R_t^{\text{terminal}} + w_I \Delta I_t + w_F \Delta F_t + w_C \Delta C_t + w_V \Delta V_t, \quad (4)$$

where ΔI_t is information gain, ΔF_t is accepted frontier progress, ΔC_t is reduction in constraint or specification uncertainty, and ΔV_t is an improvement in the ability to detect false success. The weights and units are domain-specific; this is not a universal scalar leaderboard. In mathematics, a counterexample, bridge lemma, or tightened bound can increase ΔF_t before a final theorem is proved. In engineering, a measured bottleneck or rejected design can increase ΔI_t . In system verification, a minimal trace showing that the code and specification disagree can increase both ΔC_t and ΔV_t even before either side is repaired.

Only reviewed increments enter F_t . This prevents the framework from rewarding the production of plausible intermediate text. It also gives the user a sequence of grounded decision points rather than forcing a binary choice between accepting a final answer and restarting from the original prompt.

Figure 2 contrasts a session-limited trajectory with the persistent runtime model. A session-limited agent eventually spends its budget on reconstruction and repeated work. Argus instead instantiates a recurrent Manager–Planner–Engineer–Reviewer loop inside each campaign Stage. Reviewed state persists below the loop, while Manager-controlled advance, rollback, and rejected branches change the route taken by later missions.




Figure 2. Operational model of persistent, pivot-capable long-horizon reasoning. The left card shows a session-limited route ending in reset. The central recurrent role loop is reused across eight Manager-controlled research Stages, while the blue route, orange rollback, and red rejected branches show that progress need not be monotone. Reviewed memory, Skills, tools, verifiers, routing, and task-contract refinements persist across missions and become user-visible decision state. The purple path to future SFT/RL is a research hypothesis; the current runtime changes persistent state while keeping model parameters fixed. Geometry is conceptual rather than an empirical scale.

3.4 Role-State and Stage Dynamics

The roles are not four sequential stations that execute once. Within a stage, the campaign repeatedly applies the transition relation

$$M \rightarrow P \rightarrow E \rightleftharpoons R \rightarrow M, \quad (5)$$

where M , P , E , and R denote Manager, Planner, Engineer, and Reviewer. Reviewer continuation returns work to Engineer; an accepted or blocked verdict returns control to Manager; and a Manager hold or rollback reopens planning. The relation is an explanatory state-machine model of the implemented control flow, not a claim that every transition must invoke a fresh model session.

The user does not need to approve every low-level action. Instead, the runtime raises user-visible decision points when evidence exposes an ambiguous tradeoff, contradicts a standing assumption, requires a material change to ι , or cannot determine the correct constraint from artifacts alone. This keeps repetitive execution autonomous while preserving human authority over the meaning of success. User clarification returns through Manager and becomes part of the next persisted contract rather than an ephemeral chat instruction.

Let g_n be the current stage after the n th Manager decision. Legal transitions are

$$g_{n+1} \in \{g_n, \text{next}(g_n)\} \cup \text{prev}(g_n), \quad (6)$$

corresponding to hold, advance to the immediate next stage, or rollback to an earlier stage. For the research vertical, the ordered stages are research, plan, benchmark, run, analysis, draft, review, and submission. Stage index and task-native quality need not improve on every cycle: experiments can fail, review can reject a branch, and later measurements can invalidate earlier assumptions. The intended ascent is progress in the accepted frontier over many cycles, not monotonic improvement at each transition.

3.5 Dense-Intelligence Density

The website expresses useful intelligence per unit wall-clock time as

$$\rho_I(T) = \frac{1}{T} \int_0^T \dot{N}_{\text{tok}}(t) \eta_r(t) \eta_a(t) \eta_v(t) dt. \quad (7)$$

$\dot{N}_{\text{tok}}(t)$ is the instantaneous rate of token-mediated reasoning and action. The factors $\eta_r(t)$, $\eta_a(t)$, and $\eta_v(t)$ represent the fractions converted into relevant reasoning, effective action, and valid verification. The product is intentionally stricter than token throughput: a run that repeatedly reads the same state, emits activity without changing an artifact, or accepts unchecked claims can consume many tokens while contributing little to ρ_I .

Operational efficiency factors. The three factors describe how work is distributed within an analysis window. For W , let $\mathcal{U}_k(W)$ be the attributable units for $k \in \{r, a, v\}$: observable reasoning segments, executed actions, or verification segments. We define

$$\eta_k(W) = \frac{\sum_{u \in \mathcal{U}_k(W)} w(u) q_k(u)}{\sum_{u \in \mathcal{U}_k(W)} w(u)}, \quad 0 \leq q_k(u) \leq 1, \quad (8)$$

where $w(u)$ is the attributable Token count for reasoning and verification, and either one action or its measured cost for action efficiency. The attribution label $q_k(u)$ records conversion into the intended function:

- **Reasoning efficiency η_r .** A direct unit derives, revises, or rules out a named claim, lemma, blocker, or decision. Repeated context, narrative transitions, and notation-only rewriting are transition cost.
- **Action efficiency η_a .** A direct action changes an artifact, result set, accepted decision, or explored branch. A falsification action is productive when it certifies a dead branch; retries and interrupted no-op work are tracked separately. State synchronization and packaging form the auxiliary action category.
- **Verification efficiency η_v .** A direct verification unit compares a named claim with an explicit criterion or artifact, reruns an independent check, identifies a concrete defect, or issues a scoped verdict. Schema repair, reformatting, and checkpoint canonicalization enable review but belong to verification overhead rather than direct checking.

We report a direct-work value and an inclusive value that also counts successful coordination and state maintenance. Equation 7 uses the three factors as a multiplicative bottleneck model. Section 6.5 measures them in the mathematical campaign, and Appendix D gives the full calculation.

3.6 Coupled Expert, Contract, and Runtime Evolution

The original website object separates model parameters from the operating state around the model. We extend it to expose the coupled state central to this report:

$$\begin{aligned} X_{t+1} &= F(X_t, e_t, r_t), \\ K_{t+1} &= \mathcal{A}(K_t, K'_t, e_t, r_t, u_t), \\ H_{t+1} &= U(H_t, \tau_t, E_t, K_{t+1}), \\ \theta_{t+1} &= \theta_t. \end{aligned} \quad (9)$$

F denotes the user-facing update induced by inspected evidence, alternatives, and uncertainty; it does not assume that every mission requires synchronous human input. $\mathcal{A}$ is the guarded contract-admission operator from Equation 3. U updates the persistent agent runtime using only admitted results E_t and the current contract. The coupling is bidirectional: agent work gives the user better evidence about the task, while refined user judgment changes what later agent work attempts and verifies.

with the public state definition

$$H_t = \{\text{Memory, Skills, Tools, Verifiers, Routing}\}. \quad (10)$$

E_t is the subset of results admitted from mission t . The equality for θ is a scope condition: online evolution does not require a gradient update to the underlying model. Section 4 refines H_t into named ownership surfaces and tracks task/evaluation

definitions as an additional component Q_t . The update U is partial; many missions change only memory, and some change no reusable component.

3.7 Capability-Set Expansion

Let C_t denote the effective agent capability set available at time t . The website writes review-gated admission and the desired breadth/depth behavior as

$$\begin{aligned} C_{t+1} &= C_t \cup \{c : \text{Verify}(c, E_t) \geq \varepsilon\}, \\ W_{t+1} &\geq W_t, & D_{t+1}(d) &\geq D_t(d). \end{aligned} \quad (11)$$

where W_t is frontier width across domains and $D_t(d)$ is depth within domain d . The inequalities define retention desiderata over validated capability. Skills can become stale, verifiers can be revised, and performance can regress on harder instances. A negative result can still add value by excluding a failed branch even when the effective capability set does not grow.

Capability expansion is only one half of the mechanism. The user-facing state can improve even when C_t does not expand: a verified impossibility, contradictory specification, or cost tradeoff may reduce uncertainty about which objective is worth pursuing. Conversely, a new Skill is not useful if it accelerates work under an invalid contract. The framework therefore evaluates runtime reuse and contract quality together rather than equating self-evolution with a growing tool library.

3.8 Research Value

Dense activity matters only when it produces new, valid, reusable information or improves the user’s verified understanding of the task. The website therefore defines the accumulated research value over horizon T as

$$V_R(T) = \int_0^T \rho_I(t) \Delta I(t) p_{\text{valid}}(t) \eta_{\text{reuse}}(t) dt, \quad (12)$$

where $\Delta I(t)$ is information gain, $p_{\text{valid}}(t)$ is the probability that the claimed increment survives review, and $\eta_{\text{reuse}}(t)$ is the fraction that remains useful for later work. Equation 12 separates Token volume from the properties that make a result useful: information gain, validity, and reuse. The experiments measure these components in their native task units.

3.9 The Information Advantage of Process Data

Finally, the public formalization distinguishes the accepted artifact from the trajectory that produced it:

$$D_{\text{process}} = \{(s_k, a_k, e_k, r_k, \Delta H_k)\}_{k=1}^N \supsetneq D_{\text{final}} = \{y^*\}. \quad (13)$$

The process record contains states, actions, measurements, review feedback, and runtime-state changes, including failed branches omitted from y^* . This additional information can reduce repeated search, guide later analysis, and provide grounded material for later skill construction or evaluation. Argus retains typed, privacy-preserving observations and verdicts rather than private chain-of-thought transcripts.

The set inclusion in Equation 13 can be strengthened into a decision-theoretic statement. Let P denote a typed process record and let $Y = g(P)$ be its final-artifact projection. For a downstream task q with target Z_q , loss ℓ_q , and decision policy π , define the minimum achievable risk

$$\mathcal{R}_q(X) = \inf_{\pi} \mathbb{E}[\ell_q(\pi(X), Z_q)]. \quad (14)$$

Proposition 1: process-data dominance

If $Y = g(P)$, then $\mathcal{R}_q(P) \leq \mathcal{R}_q(Y)$ for every downstream decision problem q . The inequality is strict for some q whenever two process records can yield the same final artifact while implying different optimal next actions. The proof is immediate: every policy using Y can be reproduced from P by first applying g , while the reverse simulation need not exist. This is the Blackwell-information ordering specialized to a research trajectory (Blackwell, 1953).

Dominance is informational, not computational. A raw trajectory can be too large, stale, or contradictory to use efficiently. For a context budget b , the relevant object is therefore a typed compression $\psi(P)$ that trades downstream risk against retrieval and validation cost:

$$\psi_b^* = \arg \min_{\psi: \text{size}(\psi(P)) \leq b} \mathbb{E}_{q \sim \mathcal{Q}} [\mathcal{R}_q(\psi(P)) + \lambda_c C_{\text{read}}(\psi(P))]. \quad (15)$$

Equation 15 explains why an append-only event tape and a bounded reviewed checkpoint serve different purposes. The tape preserves the more informative experiment; the checkpoint approximates a decision-useful compression under a finite context budget. Failed branches belong in the compression when they change the next optimal action, not merely because they occurred.

3.10 Reusable State and Compounding Intelligence

Process data compounds only when an admitted state update improves later work. Let $H_t \oplus \Delta H_t$ denote the state after accepting mission t , and let $q_{t+1:t+L}$ be the next L tasks. We define the discounted reuse value

$$G_L(\Delta H_t) = \sum_{j=1}^L \gamma^{j-1} [\mathcal{R}_{q_{t+j}}(H_t) - \mathcal{R}_{q_{t+j}}(H_t \oplus \Delta H_t)]. \quad (16)$$

$G_L > 0$ means that the accepted memory, skill, verifier, routing rule, or certified dead branch reduces future loss under a prespecified task distribution. $G_L < 0$ captures negative transfer. This definition turns “compounding intelligence” from monotone rhetoric into a counterfactual claim: reuse must outperform a frozen state on matched future tasks. Workflow-memory experiments show that induced reusable routines can improve success and reduce steps (Wang et al., 2024b). Section 7 connects this quantity to the current Argus traces.

The website’s density and research-value equations can then be joined into a protocol-specific *verified reusable yield*:

$$\mathcal{Y}_{\text{PRI}}(W) = \frac{\sum_{t \in W} p_{\text{valid},t} [\Delta I_t + \lambda_g G_L(\Delta H_t)]}{\sum_{t \in W} N_{\text{tok},t}}. \quad (17)$$

ΔI_t is immediate task-native information gain, such as a score improvement, theorem strengthening, or certified branch elimination; $p_{\text{valid},t}$ is estimated by an external evaluator or a calibrated review procedure; and G_L measures future reuse. λ_g converts immediate and future value into one protocol-specific unit. Equation 17 connects immediate research output to future reuse. The current substitutions are reported in Appendix D.

3.11 Review as Selective Error Correction

Let C denote whether a proposed state increment is correct, A whether the Reviewer accepts it, $p = \Pr(C = 1)$ the proposal base rate, $\alpha = \Pr(A = 1 \mid C = 1)$ Reviewer sensitivity, and $\beta = \Pr(A = 1 \mid C = 0)$ false-acceptance rate. Bayes’ rule gives

$$\Pr(C = 1 \mid A = 1) = \frac{\alpha p}{\alpha p + \beta(1 - p)}. \quad (18)$$

If $\alpha > \beta$, accepted state is more precise than the proposal stream; review acts as a selective error-correction channel. Its operating point balances accepted precision, recall, and Token cost. Section 6 reports the observed repair and recovery rates.

3.12 Role-Resolved Vertical Traces

Endpoint metrics do not show how an Agent system actually conducts research. We therefore project each mission trajectory in Equation 2 onto the four role-owned actions

$$\phi(\tau_t) = (m_t, p_t, x_t, r_t, \Delta H_t), \quad (19)$$

where m_t is the Manager’s objective or stage decision, p_t is the Planner’s bounded task and dependency decision, x_t is the Engineer’s artifact-producing execution, r_t is the independent Reviewer’s structured completion decision, and ΔH_t is the durable state update exposed to later missions. The decomposition makes two authority constraints explicit: only the Manager changes campaign stage, while mission completion requires a fresh Reviewer verdict after every normal Engineer round. Historical self-review labels remain in older traces but do not describe the current production chain; the Planner may propose future work but cannot certify the mission.

A vertical trace is useful when these tuples remain linked across many missions. It shows whether the objective survived failures, whether the Planner consumed previously accepted state, whether the Engineer changed real artifacts, whether the Reviewer rejected or redirected incomplete work, and which decisions became inputs to the next mission. This explanatory model organizes the Agent workflow; Section 6.5 applies it to one mathematical campaign.

4 Argus Runtime

4.1 System Overview

Argus is a verification-centered runtime built around a user, four model-driven roles, and three system planes. The user supplies standing intent, domain judgment, and authority over material tradeoffs. The *control plane* turns that intent into a durable task contract and schedules work; the *execution plane* performs one bounded mission against real tools and artifacts; and the *record plane* stores evidence, decisions, and contract history without deciding whether the work is scientifically complete. The distinction is operational: control owns scheduling and contract transitions, execution owns work and mission review, and the record plane owns the immutable record but no completion decision.

The four roles divide research authority as summarized in Table 1. The Manager spans the campaign rather than one execution round. The Manager is also the user-facing boundary: it preserves standing intent, surfaces consequential ambiguity, and

commits authorized contract refinements. The Planner authors future work and the Engineer performs it. Every normal mission round then passes to a fresh independent Reviewer; historical Engineer-self-review fields remain only for compatibility with older traces. The interfaces are structured objects rather than untyped prose, which makes ownership and recovery behavior explicit.

The notation K_t is a report-level abstraction over persisted campaign identity, objectives, Stage state, constraints, verifier obligations, and decision records. The current implementation distributes these fields across runtime surfaces rather than storing one monolithic object named K_t .

Role	Primary responsibility	Authoritative output
Manager	Preserve standing intent; commit task contract, lifetime, route, and campaign stage	Contract refinement, user escalation, campaign division, and stage transition
Planner	Decompose the current verified frontier into bounded tasks and dependencies	Task specifications, plan status, verifier obligations, and stage checklist updates
Engineer	Modify artifacts, invoke tools, and run experiments for one round	Execution result, artifacts, measurements, and checkpoint handoff
Reviewer	Independently inspect every normal round against the current contract	done, continue, blocked, or replan_requested; verified frontier update; contradiction or expert-question report

Table 1. Authority separation in Argus. The Engineer executes; a fresh independent Reviewer controls admission after every normal mission round; Manager owns contract and Stage transitions at the user-facing boundary.

Algorithm 1: Reviewed mission loop

Input: standing intent t , contract K_t , user state X_t , runtime state H_t , backlog B , budget b
 $c \leftarrow \text{ManagerCommit}(t, K_t)$; persist campaign identity and contract
while b remains and c is not complete **do**
 $q \leftarrow \text{Planner}(B, H_t, c)$; claim one bounded task
 $\tau_t \leftarrow []$
 repeat
 $(y, e) \leftarrow \text{Engineer}(q, H_t)$
 $r \leftarrow \text{Reviewer}(K_t, q, y, e)$
 append (q, y, e, r) to τ_t
 if $r = \text{continue}$ **then** $q \leftarrow r.\text{next_action}$
 until $r \in \{\text{done}, \text{blocked}, \text{paused}, \text{replan_requested}\}$
 $E_t \leftarrow \text{AdmitResult}(e, r)$
 $X_{t+1} \leftarrow \text{SurfaceEvidence}(X_t, E_t, r)$
 $K_{t+1} \leftarrow K_t$
 if $r = \text{replan_requested}$ or contract refinement is proposed **then**
 $K_{t+1} \leftarrow \text{ManagerAdmit}(K_t, E_t, r, u_t)$
 $H_{t+1} \leftarrow U(H_t, \tau_t, E_t, K_{t+1})$; persist outcome, trace, contract, and state
 $t \leftarrow t + 1$; refill B when required
end while
Output: accepted frontier, inspectable trajectory, refined task contract, and updated expert/runtime state

4.2 Bounded Missions over Durable Contracts

The long-lived object in Argus is the campaign, not a provider transcript. A campaign has a persisted identity, standing intent, current task contract, accepted frontier, and explicit unresolved questions. Its work is divided into missions with explicit outcomes. The scheduler assigns one mission at a time, records its result, and advances only at a clean mission boundary. Mission assignment is transactional, preventing duplicate work under concurrency, and resumed execution is tied to the persistent campaign identity.

Engineer and Reviewer calls use fresh provider sessions for each round. Cross-session continuity comes from an ordinary shared CHECKPOINT.md containing durable state, evidence references, open questions, and the next step. This follows the broader move from monolithic context toward managed memory tiers and reusable workflows (Packer et al., 2023; Wang et al., 2024b; Xu et al., 2025). The Engineer updates the checkpoint after execution. The Reviewer then reads and corrects the same file and is its final editor for that round. The full history remains in artifacts and the event record.

This design makes restart and upgrade behavior part of the method. A running process hands off only at a mission boundary,

and the replacement resumes from the same persistent campaign identity. After an interruption, replay or reassignment begins from committed campaign state rather than reconstructing the task from a model transcript.

The same persistence supports user interaction at a larger timescale. The user does not need to repeat the original context or monitor every tool call. When a Reviewer finds that the current verifier cannot distinguish alternatives, that a constraint conflicts with observed reality, or that multiple valid tradeoffs remain, the issue is surfaced with evidence and options. The resulting clarification is committed through Manager into the durable contract. Later missions therefore begin from the refined understanding rather than from either the original ambiguous prompt or an agent-invented replacement goal.

4.3 Autonomy between User Decision Boundaries

The runtime does not require a human to supervise each round. The harness assigns local authority to the four roles: Manager maintains the campaign and contract, Planner creates bounded executable work, Engineer acts on real artifacts, and an independent Reviewer decides whether evidence supports completion, continuation, blocking, or replanning. Checkpoints, event history, task claims, external-job state, and mission outcomes survive session replacement. Consequently, the runtime can execute for long intervals without synchronous human input and surface a user decision only when evidence exposes a material value tradeoff, missing information, or operator-owned constraint.

This is stronger than adding a human-approval button. The autonomous path must remain capable of selecting routes, running experiments, rejecting invalid work, consolidating intermediate progress, and ending an unproductive mission. Human input refines the campaign at consequential boundaries; it is not the mechanism that keeps the ordinary loop alive.

4.4 Verification as the Pivot Gate

All role invocations pass through a common instrumentation layer that records usage and associates each call with the mission trace. The Engineer and Reviewer share the same artifact state, allowing review of the actual outputs and execution record rather than only the Engineer’s summary. A typed, append-only trace is the canonical timeline; user interfaces are projections over that record.

The source of a mission completion verdict is explicit: a fresh Reviewer returns a structured verdict after every normal Engineer round. Review evaluates more than terminal completeness. It can admit a bounded frontier increment, reject an unsupported claim, identify a faulty verifier or inconsistent specification, request another experiment, or ask Manager to replan and surface a user decision. Historical `review_source=engineer_self_review` values describe earlier runtime versions and are retained only so longitudinal traces remain interpretable. Completion never comes from filenames, Token volume, content hashes, or prose keywords. The corresponding artifacts and execution records remain inspectable, while credential redaction preserves the evidence boundary.

Verification is therefore the gate between persistence and pivoting. Agent execution externalizes hypotheses into artifacts and measurements; Reviewer turns those outputs into calibrated claims and unresolved alternatives; Manager presents material choices in terms the user can inspect; and user judgment updates the contract governing subsequent execution. Without verification, persistence can compound error. Without persistent self-evolution, verification repeatedly finds the same errors. Their combination converts a long trajectory into cumulative problem understanding and reusable capability.

4.5 Staged Approach Self-Evolution and Contract Pivoting

Equation 9 separates the model from the state around it. Table 2 identifies how each component can change and who owns the committed update. The ownership model is intentionally not uniform. Memory and skills use a work-versus-certification split: the Engineer or Scientist produces a candidate, and the Reviewer commits the retained form. Tools and procedures are system-configured. Stage checklists are Planner-owned, with Reviewer feedback rather than a second commit gate. Routing policy is Manager-committed, while task definitions are Planner-authored and scheduler-committed. The task contract is different from all of these: evidence and Reviewer findings may propose a change, Manager commits operational refinements, and the user retains authority over standing intent and value-laden tradeoffs that evidence alone cannot settle.

We use verifier-gated runtime self-evolution to mean the staged evolution of the approach used to move from the current objective toward satisfaction of the current constraints. The underlying model parameters remain fixed, while reviewed missions change the persistent search state available to later work. Each Stage maintains an approach record Π_t containing route hypotheses, actions, task-native evidence, Reviewer verdicts, retained failures, and the next admissible frontier. A complete update cycle has four parts: (1) an execution trajectory produces a candidate memory, skill, procedure, verification rule, routing decision, or task definition; (2) the responsible role checks the candidate against artifacts and task-native evidence; (3) the authorized owner commits, revises, or rejects the update; and (4) a later mission retrieves the retained state as part of its starting context or execution policy. Activity that does not survive this commit-and-reuse path is not counted as self-evolution. Contract pivoting is distinct: it is the evidence-backed refinement of K_t and X_t , with user clarification or authorization where required. Neither mechanism implies that every mission improves the system or changes the contract.

This preserves a strong self-evolution claim without conflating it with goal mutation. K_t defines the current objective, constraints, and verifier; Π_t records how the runtime has learned to navigate inside that envelope. If a verified contract

refinement changes the envelope, the prior approach record remains staged evidence of which routes failed and which procedures or verifiers remain reusable.

State	Change source	Commit owner	Persistent form
<i>M</i> : memory	Engineer trajectory	Reviewer	Curated checkpoint and event-derived journal
<i>S</i> : skills	Engineer / Scientist	Reviewer	Versioned skill library
<i>A</i> : tools and procedures	System configuration	System configuration	Versioned tool and procedure registry
<i>V</i> : verification	Planner	Planner	Stage checklist; Reviewer supplies feedback
<i>R</i> : routing and roles	Runtime policy	Manager	Campaign routing policy
<i>Q</i> : tasks and evaluations	Planner	Scheduler	Backlog task specifications and scopes
<i>Π</i> : staged approach	Mission trajectories and Reviewer verdicts	Manager / Planner	Accepted frontier, rejected routes, route rationale, and next approach
<i>K</i> : task contract	User clarification, evidence, and Reviewer findings	Manager / user	Versioned intent, objective, constraints, verifiers, and provenance

Table 2. Attributable ownership of persistent and revisable state. A mission usually updates only a subset of these components; material changes to standing user intent are not delegated to the runtime.

Two knowledge surfaces carry reusable experience. Versioned skills store procedures that can be matched to later tasks. A project knowledge base stores source-linked records and synthesized pages derived from reviewed outcomes. Neither surface is treated as automatically correct: entries can be revised, archived, or retired when later results contradict them.

This mechanism can improve a later mission without changing the model itself. A retained failure can prevent a repeated dead end; a reviewed procedure can shorten environment inspection; a verifier can reject a previously accepted shortcut; and a routing update can assign independent review to a higher-risk task. The mechanism does not imply monotonic improvement. Some missions commit no reusable state, retained state can become stale, and a harder task distribution can increase cost even after useful state has accumulated.

The user-facing benefit is complementary. A retained failure can show that a requested optimization is physically or economically irrelevant; a stronger verifier can expose that the original success metric was gameable; and a proof counterexample can replace a vague conjecture with a sharper reachable frontier. The runtime is not learning the user’s preferences covertly. It is making the evidence and tradeoffs explicit enough for those preferences to be stated and revised deliberately.

4.6 Reliability and Resource Governance

Long-running systems fail through operational drift even when individual model calls are strong. Argus therefore combines round-level progress classification with mission-boundary replanning. Work that repeatedly produces no decision or artifact change is stopped or reformulated, while long-running external jobs are tracked separately from model reasoning. These mechanisms bound execution; scientific correctness remains with the task evaluator and independent Reviewer.

Short Planner DAG nodes expose a subtle control problem: their default three-round budget is much smaller than the default absolute thresholds for semantic stall, soft escalation, and hard escalation. The current runtime fits each enabled guard to the reachable range when `max_rounds` shrinks. For budget m , semantic stall and soft escalation must satisfy $g \leq m - 1$, while hard escalation may fire on the final round and therefore satisfies $g \leq m$. Explicit zero still disables a guard, and the default 500-round configuration is unchanged. A two-round budget does not silently turn a configured multi-round stall detector into a one-strike policy: the impossible detector is disabled unless the caller explicitly requests threshold one. This coordination preserves meaningful stop and replan signals instead of falling through silently to the generic maximum-round outcome.

These liveness guards do not decide scientific truth and do not freeze a mistaken contract. They create bounded points at which Reviewer and Planner can distinguish three cases: the current route is unproductive, the evidence frontier has advanced but terminal work remains, or the task contract itself needs user-visible refinement. This distinction is important because “not done” is not synonymous with “no progress,” and “cannot verify” is not synonymous with “incorrect.”

We call this property *honest stopping*: the harness can end the current unit of work without converting uncertainty, exhaustion, or an external dependency into a success claim. The implementation keeps completion, incomplete research, stalling, blocking, infrastructure blocking, failure, and operator interruption as separate outcome dimensions. Reviewer may return `blocked` or `replan_requested`; a hard escalation ends the mission with a planner-readable reason. Semantic stall accumulates only from the Reviewer’s explicit structured `forward_progress=false`; a missing or malformed signal is treated as unknown rather than inferred from prose or filesystem activity. Stopping therefore preserves both epistemic honesty and the information needed to resume, decompose, or change approach.

Resource accounting is centralized at model-call and external-job boundaries. Budgets are checked before work begins and reconciled after completion, preventing individual roles from expanding their own allocation. Concrete scheduling, sandboxing, and deployment mechanisms are implementation choices rather than part of the scientific claim.

5 Empirical Methodology

We evaluate Argus across seven benchmark arenas spanning software repair, GPU-kernel optimization, language-model training, training-speed optimization, research-assistant tasks, and data synthesis. Each benchmark is reported in its native unit. SWE-Bench Pro is one benchmark in this suite; its sequential task log also enables deeper analyses of runtime evolution and Reviewer intervention. The evaluation does not collapse persistence and pivoting into a single score. Endpoint benchmarks measure whether the runtime can still deliver; longitudinal traces measure intermediate frontier updates, verification-driven correction, retained state, and evidence-backed reframing.

5.1 Research Questions

RQ1: Benchmark performance. What task-native outcomes does Argus achieve across the seven benchmark arenas, and how do they compare with the reference reported for each arena?

RQ2: Fixed-model agent-side evolution. Within the SWE-Bench Pro run, how does solve-time resource demand change as reviewed Skill, Wiki, verification, and routing state accumulates while the underlying model parameters remain fixed?

RQ3: Reviewer routing and recovery. Within the same SWE-Bench Pro run, how often is an independent Reviewer invoked, and how many initially rejected solutions are recovered after Reviewer feedback?

RQ4: Verified intermediate progress. Do the retained mathematical and paper-production traces show useful admitted progress before terminal completion, including certified dead routes, stronger bounds, corrected constraints, and verifier-driven rollback?

RQ5: Evidence for verified pivoting. How do retained trajectories expose changes in the operational objective, constraints, or acceptance criteria, and which claims about user benefit remain unmeasured without a prospective interaction study?

5.2 Benchmark Suite

SWE-Bench Pro evaluates repository-level software-engineering tasks with executable acceptance tests (Deng et al., 2025). The remaining arenas evaluate B200 kernel optimization, nanochat training on B200 and H100, the nanoGPT speedrun, AARRI-Bench, and Math-Reasoning Data Synthesis from the Arbor suite (Lin et al., 2026; Karpathy, 2025, 2026; Jordan and contributors, 2024; Wang et al., 2026; Jin et al., 2026). Because rank, bits per byte, elapsed time, solve rate, and pass-gap are incompatible metrics, we do not average them into one score.

Benchmark	What the agent must do	Evaluation metric	Better
SWE-Bench Pro	Repair a real repository issue by editing production code; the submitted patch must satisfy task-specific executable acceptance tests.	Fraction of tasks resolved by the official verifier	Higher
SOL-ExecBench	Optimize real GPU kernels without changing their outputs, targeting execution close to an analytically derived hardware speed-of-light bound.	SOL score, global rank, and high placements after correctness checks	Higher
nanochat B200	Improve the nanochat training program under a fixed five-minute run on one B200 GPU.	Validation bits per byte (BPB), a compression-oriented language-model metric	Lower
nanochat H100	Solve the same five-minute nanochat training objective on one H100; it is a separate hardware-specific optimization problem.	Validation BPB under the H100 protocol	Lower
nanoGPT speedrun	Modify the training system so an eight-H100 run reaches validation loss 3.28 as quickly as possible.	Mean time to the target over ten verifier runs	Lower
AARRI-Bench	Complete 82 granular tasks representing the judgment, diligence, and professional practice expected from a research intern.	Number and percentage of tasks solved	Higher
Math-Reasoning Data	Build a pipeline that generates AIME-style reasoning problems that are difficult on the first attempt but solvable with additional attempts.	Mean pass@4 – pass@1 gap	Higher

Table 3. Definitions of the seven benchmark arenas. B200 and H100 nanochat runs are kept separate because hardware changes the optimization problem.

The six non-SWE arenas use GPT-5.5 through Codex. The SWE-Bench Pro evaluation uses GPT-5.5/xhigh through Copilot for both Direct Copilot and Argus. The full 731-task comparison reports approximate accuracy and the aggregate Token ratio

$$R_{\text{tok}} = \frac{N_{\text{Argus}}^{\text{total}}}{N_{\text{Copilot}}^{\text{total}}} \approx 1.41. \quad (20)$$

Raw Copilot Token totals and per-Wave resource traces were not retained; therefore the longitudinal analysis is Argus-only.

5.3 SWE-Bench Pro Runtime Self-Evolution Analysis

For a completed Argus Wave w with task set $\mathcal{T}_w$, we measure

$$\bar{I}_w^{\text{solve}} = \frac{1}{|\mathcal{T}_w|} \sum_{i \in \mathcal{T}_w} I_i^{\text{solve}}, \quad (21)$$

$$\bar{T}_w^{\text{agent}} = \frac{1}{|\mathcal{T}_w|} \sum_{i \in \mathcal{T}_w} T_i^{\text{agent}}. \quad (22)$$

I_i^{solve} contains all model calls in the accepted task trajectory, including routing, Skill adaptation, execution, and any independent review. T_i^{agent} measures active workflow time. Orchestration wait, environment preparation, external verification, infrastructure recovery, and post-task knowledge maintenance are excluded.

Completed Waves are summarized by task-weighted windows: W1–6 (startup), W7–12 (early reuse), W13–18 (composition shift), W19–22 (mature operation), and W23–24 (late difficult tasks). Two incomplete Waves are omitted from the grouped means. The primary comparison is startup versus mature operation; composition and difficulty variation remain visible.

The unit of analysis is the observed sequential window, not a matched task pair. The startup window begins with less project-specific state, whereas later windows can retrieve reviewed repository knowledge, reusable procedures, prior failure records, verification guidance, and updated routing decisions. Lower Token or time demand in a later window is therefore consistent with runtime self-evolution when the retained state removes repeated inspection or failed work. It is not, by itself, a causal estimate of the value of that state. Task identity, repository mix, execution latency, and difficulty also change across Waves, and no matched frozen-state replay is available.

5.4 SWE-Bench Pro Reviewer Analysis

A task is *Reviewer-invoked* when its historical SWE-Bench Pro trajectory contains an independent Reviewer model call; otherwise the trace carries the then-active Engineer self-review label. This split is an observed property of the archived evaluation runtime, not the current all-round independent-review policy. A Reviewer verdict of `continue` is a revision request. We report both subsequent official verifier success and the stricter case in which a later Reviewer verdict is `done`. The latter is called a *strict Reviewer rescue*.

Reviewer routing is adaptive rather than randomized. The analysis reports routing, revision, and recovery behavior over the observed task sequence.

The 41-artifact research portfolio is reported separately as a measure of research program breadth.

5.5 Representative Vertical-Trace Protocol

To complement endpoint and software-trajectory measurements, we report one Erdős–Gyárfás mathematical campaign as a representative vertical trace. The trace is role-led rather than theorem-led: it records what the Manager, Planner, Engineer, and Reviewer did during recovery, problem selection, route testing, proof production, revision, acceptance, and state retention. The reported record contains one accepted route falsification and six proof-backed research deltas; detailed mathematical statements remain in the supplementary claim table.

The protocol tracks whether the runtime retains a falsified branch, carries named blockers across missions, and admits later work after artifact-based review. The role trace, claim table, and figure provenance are released with the supplementary materials.

5.6 Theory-to-Measurement Mapping

We map the theory to four reported quantities. Immediate information gain ΔI_i is represented only in task-native units: benchmark improvement, accepted theorem delta, or certified branch elimination. Review correction is represented by Reviewer `continue`–revise trajectories and later external-verifier outcomes. Reuse value G_L is approximated observationally by later Token/time demand as Skill and Wiki state accumulates. Token conversion is represented by the role- and mission-level attribution in the mathematical case. Appendix D records the corresponding substitutions in their native units.

6 Results

The results answer two complementary questions. The seven benchmark rows test whether a verification-heavy persistent runtime can produce strong terminal outcomes. The longitudinal traces test why one-shot completion is an incomplete model: revisions recover failed solutions, mathematical intermediate states change the next theorem task, and paper campaigns alter claims when evidence invalidates their initial positive hypotheses. The latter are process observations, not a controlled human-subject demonstration of expert learning.

6.1 Benchmark Results

Table 4 is the primary empirical summary and reports all seven benchmarks in one comparison; Table 3 defines the task and metric for each row.

Benchmark	Backbone	Backend	Protocol	Argus result	Comparison
SWE-Bench Pro	GPT-5.5/xhigh	Copilot	731 software-repair tasks	$\approx 78\%$ accuracy	Direct Copilot $\approx 59\%$; $1.41\times$ aggregate Tokens
SOL-ExecBench	GPT-5.5	Codex	B200; 101 kernels	Global #6; $2\times\#1$; 7 top-3	Two head-to-head wins over Recursive
nanochat B200	GPT-5.5	Codex	5 min; $1\times\text{B200}$	0.9636 BPB	Human SOTA 0.9646
nanochat H100	GPT-5.5	Codex	5 min; $1\times\text{H100}$	0.9855 BPB	Human SOTA 0.9879
nanoGPT speedrun	GPT-5.5	Codex	$8\times\text{H100}$; $N=10$	79.77 s	Same-device human 80.18 s
AARRI-Bench	GPT-5.5	Codex	82 research-intern tasks	63/82 (76.8%)	Paper best 68.3%
Math-Reasoning Data	GPT-5.5	Codex	AIME-style synthesis	28.0 gap	Arbor 20.83; Claude 8.33; Codex 6.25

Table 4. Results across seven benchmark arenas. Values remain in task-native units and are not cross-normalized. Backbone denotes the model driving the research agent; Backend denotes the execution surface.

As a small upstream-adoption example, an Argus-optimized TileLang RWKV6 kernel was reviewed by a Moonshot AI-affiliated FLA collaborator and merged into `fla-org:main` (PR #1045; commit `c70f11c`); Appendix E records the technical details.

A separate chip-design project record supplied with this revision reports ACE-1 (Argus Compute Engine), a batch-one decode NPU prototype with synthesizable RTL, a Sky130 gate-level netlist produced through Yosys and ABC, a 128-bit DMA interface, an eight-bank 256 KiB scratchpad, and INT4 compute. Its project score is reported as 75/80 against the stated area, power, and throughput comparison. This broadens the artifact surface from software kernels to digital design, but it has a different evidence level from the merged FLA contribution: the report bundle does not contain the ACE-1 raw score inputs or an independent reproduction. We therefore treat it as an operator-reported prototype rather than an eighth benchmark row.

The table shows one runtime reaching competitive outcomes across heterogeneous tasks. The remainder of this section analyzes the SWE-Bench Pro row because it provides a long sequential trajectory and Reviewer decisions.

6.2 SWE-Bench Pro Analysis: Fixed-Model Runtime Self-Evolution

Figure 3 shows the Argus-only longitudinal trajectory. Here, self-evolution refers to changes in persistent runtime state, not to online training of the underlying model. Across Waves, accepted trajectories can update repository knowledge, reusable Skills, verification guidance, task definitions, and routing policy. These objects alter what later missions retrieve, which checks they run, and where they spend independent review.

Mature operation W19–22 uses 21% fewer solve input Tokens and 15% less active workflow time per task than startup W1–6. This pattern is consistent with accumulated state reducing repeated inspection, rediscovery, or failed execution. The lowest-Token window, W13–18, reaches a 50% reduction from startup but requires more active time, showing that Token efficiency and execution latency do not move together. The final difficult-task window rebounds in both metrics.



Figure 3. Outcome, review, and longitudinal efficiency in the 731-task SWE-Bench Pro run. Panel (a) reports the full-suite accuracy and aggregate-Token comparison, adaptive Reviewer routing, and observed review outcomes. Panels (b) and (c) aggregate Argus solve input Tokens and active workflow time over approximately six-Wave windows using task-weighted means. Relative to W1-6, W19-22 uses 21% fewer solve input Tokens and 15% less active time per task. Two incomplete Waves are omitted, while W23-24 remains visible as late difficult-task stress. Copilot per-Wave resource traces were not retained, and the window comparison is observational rather than a controlled learning ablation.

The trajectories show a lower-cost mature operating window together with visible task-composition and late-wave difficulty effects. The rebound in W23-24 prevents the startup-to-mature comparison from being read as monotonic improvement. Because the task sequence is not replayed against a frozen runtime state, the evidence supports an observational claim about system-state accumulation rather than a causal claim that any individual memory or Skill update produced the reduction.

6.3 SWE-Bench Pro Analysis: Reviewer Routing and Recovery

In the archived SWE-Bench Pro runtime, of 731 tasks, 466 (63.7%) invoke an independent Reviewer, while 265 (36.3%) carry the historical Engineer self-review route. The current production runtime reviews every normal round independently. The Reviewer requests another implementation round on 43 tasks. After revision, 34 pass the official verifier and 22 complete the strict continue→revision→done loop.

Reviewer-routed tasks consume $2.75\times$ as many solve input tokens and $1.80\times$ as much active time as historically self-reviewed tasks on average, indicating that the archived routing policy selects a harder workload. The recovery funnel is therefore more interpretable than a raw group-accuracy comparison.

Figure 4 separates the routing population from the revision-recovery path. The accepted and revise branches partition outcomes after Reviewer invocation together with 35 blocked cases; verifier pass and strict rescue are nested recovery milestones among the 43 revision-requested tasks.




Figure 4. Adaptive Reviewer routing and revision recovery. Panel (a) shows that 466 of 731 tasks invoke an independent Reviewer, while 265 use Engineer the historical self-review route; routed tasks consume $2.75\times$ the solve input Tokens and $1.80\times$ the active time on average. Panel (b) follows the routed workload: 388 tasks are accepted on first review, 43 receive a revision request, 34 of those later pass the official verifier, and 22 complete the strict review-loop rescue. Routing is task-dependent, so the cost contrast is descriptive rather than a randomized causal estimate.

6.4 Measured Process Quantities

The process theory is instantiated from two complementary traces. In the mathematical campaign, direct reasoning, verification, and action each account for approximately 56% of their respective attributed units. In SWE-Bench Pro, Reviewer revision requests lead to 79.1% official-verifier recovery and 51.2% strict review-loop rescue, while the mature window uses 21% fewer solve input Tokens and 15% less active time than startup. Appendix D gives the complete substitutions for the density, reuse, yield, and compression quantities.

The public research inventory additionally contains 41 de-duplicated artifacts across six programs, illustrating the range of research programs executed by the same runtime.

6.5 Representative Vertical Agent Trace

Figure 5 follows the retained mathematical claim frontier inside one real campaign. The case begins with an open-ended standing intent rather than a fully specified theorem-production recipe. Manager preserves that intent through budget and process failures; Planner repeatedly re-expresses the reachable objective as bounded work; Engineer retrieves sources, runs tools, and writes research artifacts; and Reviewer decides which intermediate claims, counterexamples, and proof obligations can enter durable state.




Figure 5. Reviewed claim frontier in the representative mathematical campaign. Seven retained updates progress from a falsified route (C5), through two independently re-derived baselines (C20–C21), to four campaign-frontier results (C22–C25), including the 27-to-25 strengthening and two bridge lemmas. The middle cards expose the corresponding artifacts, while the bottom bars separate direct frontier-changing work from an inclusive measure that also credits validation and state consolidation. Check marks indicate internal Argus review, not external peer review or a novelty claim.

The trace makes contract refinement concrete. Manager recovers the campaign and owns stage transitions. Planner first selects a cheap falsification test, then changes the operational success contract from “collect observations” to “produce a proof,” and later requires new missions to advance the retained result rather than restart from an easier fact. These changes preserve the standing intent while sharpening what counts as progress. Engineer performs the domain work—source retrieval, executable checks, proof writing, and result packaging. The Reviewer rejects an overstated route, requests the missing checks, and certifies the bounded theorem recorded by the artifacts.

In the reported trace, the campaign retains one falsified route and six proof-backed deltas, including one stricter bound and two bridge results. None needs to be the final proof to carry reward under Equation 4: each changes the accepted frontier or excludes future work. Accepted work changes the next Planner input, rejected work remains available as a dead branch, and open blockers become explicit future tasks. The trace therefore captures an evolving problem understanding and agent state rather than a sequence of isolated answers.

6.6 Efficiency Attribution in the Mathematical Case

We apply Equation 8 to the theorem-production window covering Rounds 12–17. The analysis contains 18 bounded missions: six changed the theorem frontier, four directly checked or re-certified mathematical results, seven consolidated review state and documentation, and one overlap-analysis mission was interrupted without an admitted result. Table 5 uses mission purpose as a coarse attribution rule for process allocation.

Efficiency view	Direct	Auxiliary	Failed/no-op	Unit and interpretation
Reasoning	56.0%	39.1%	4.9%	Engineer reasoning-output Tokens; direct means a frontier-changing proof mission.
Verification	55.4%	35.0%	9.6%	Reviewer reasoning-output Tokens; direct means proof comparison or re-certification.
Action	55.6%	38.9%	5.6%	Mission count; direct comprises six frontier missions and four verification missions.

Table 5. Efficiency attribution for the representative mathematical case. Auxiliary work includes result packaging, checkpoint updates, and documentation consolidation.

The normalized interpretation is concrete. Of every 100 Engineer reasoning-output Tokens, approximately 56 occurred in frontier-changing proof missions, 39 in verification or state-maintenance work, and 5 in the interrupted branch. Of every 100 Reviewer reasoning-output Tokens, approximately 55 directly checked a proof package, 35 maintained formal review state, and 10 were spent in the interrupted analysis. The strict action score is $10/18 = 55.6\%$; if the seven successful auxiliary missions are also credited, the inclusive score is $17/18 = 94.4\%$. Cost weighting makes the interruption more visible: the strict and inclusive action scores become 56.1% and 89.7%, respectively.

The trace also clarifies what the numerator should reward. The early witness check that killed a proposed reduction counts as an efficient action because it pruned a false route before a long proof attempt. In Round 15, selecting one internal triple, deriving the degree bounds, and closing the counting inequality are direct reasoning; rewriting the result into the structured claim record is auxiliary work. Reviewer inspection of those steps is direct verification, whereas updating the working checkpoint is coordination overhead. This separation prevents activity volume from being mistaken for research progress while preserving necessary coordination as a visible cost.

6.7 Long-Horizon Paper-Production Case Study

The benchmark suite evaluates task endpoints, while the mathematical trace follows one research objective in depth. A complementary question is whether the same runtime can preserve evidence and refine the research contract across a complete paper-production lifecycle. We therefore reconstruct six projects spanning evaluation reliability, vision-language matching, test-time adaptation, GUI agents, multimodal hallucination, and model quantization. The initial objectives and compute environments were provided externally; the retained Argus records identify the Manager, Planner, Engineer, and Reviewer as the actors responsible for Stage transitions, experiments, manuscript construction, review, and submission checks.

Figure 6 summarizes the portfolio. All 6 of 6 canonical pipelines reach the final submission Stage. In aggregate, they span 640 campaign-hours, 254 bounded missions, 576 Engineer rounds, 286 Reviewer revision verdicts, 89 session rolls, and 16 Manager rollbacks. Campaign-hours are summed across projects and are not calendar time because several projects overlap.




Figure 6. Scientific outcomes and production load across six paper campaigns. Each card pairs the final manuscript’s first page with its domain, scoped research question, and principal task-native result, including diagnostic and negative findings. The bottom strip reports the shared production process: six completed pipelines, 640 aggregate campaign-hours, 576 Engineer rounds, 286 Reviewer revisions, 89 session rolls, and 16 Stage rollbacks. Manuscript completion does not imply venue acceptance, novelty, or external peer review.

The trajectories are inconsistent with a one-pass account of either autonomous or human-directed paper writing. Each manuscript states a scoped research question and a task-native outcome rather than merely presenting a generated document. The

portfolio includes uncertainty-aware leaderboard reporting, a failed compositional gate, an over-restrictive adaptation controller, a narrow GUI reliability gain, a failure-mode audit of multimodal wrappers, and a negative result for static quantization proxies. Several projects therefore become diagnostic or negative results after their original positive hypothesis fails. Reviewer-gated state makes that scientific pivot explicit instead of silently replacing the failed branch. The persistent invariant is the standing intent to produce an evidence-grounded scientific contribution; the operational claim is allowed to narrow when the positive method hypothesis does not survive verification.

Figure 7 resolves the multimodal-hallucination campaign in greater detail. During a dense 12-hour search window, seven rollback decisions reject baseline coverage or proposed mechanisms that are unreproduced, base-identical, or missing their preregistered signal. Planner then changes the claim from a positive mitigation method to a diagnostic negative-results audit. This is the report’s clearest example of refinement rather than goal drift: seven measured no-go decisions justify the change, the failed routes remain in the record, and the new claim is harder to overstate because it is bound to the canonical matrix. Engineer completes a five-method by three-benchmark matrix with 4,500 official-scored rows; Reviewer binds the resulting no-op and degradation claims to those outputs. Two later submission checks return the project to benchmark to repair scorer provenance and GPU telemetry before the final 10-page AAAI-formatted package completes the pipeline.




Figure 7. Representative 163.6-hour paper-production trajectory. Panel (a) plots the Manager-controlled Stage against campaign time: seven early no-go decisions prune weak method routes, the project pivots from a positive method claim to an audit, and two late rollbacks repair the submission evidence. Panel (b) reduces the trace to five scientific episodes—prune, pivot, experiment, write, and repair—while retaining the 15-cell, 4,500-row experiment and the final 10-page manuscript.

The process layer is similarly nontrivial. Even the shortest case requires 28 Engineer rounds, 19 Reviewer revisions, and 13 session rolls over 20.8 hours. The two AAAI-formatted cases finish as 10-page anonymous submissions, while the four ACL-formatted cases contain 10–13 pages. Across all six projects, the retained academic, layout, and infrastructure histories contain 436 automated review snapshots. These are process records rather than external assessments of paper quality.

The case study also exposes a consistency failure. The compositional-matching pipeline reaches a complete submission Stage and produces its final PDF, while the last stored assurance object remains `BLOCKED` on a Manager-stage authority check. This discrepancy does not change the paper artifact, but it shows that derived certification state can lag the canonical pipeline state. Public release automation should reconcile those two views before presenting assurance as a final verdict.

7 Analysis and Discussion

7.1 The Core Distinction: Completing a Task versus Discovering It

Most agent evaluations present a fixed prompt, environment, and success function. That abstraction is necessary for benchmarking but incomplete for open-ended reasoning. A long-horizon campaign can fail even when every action is competently executed if the initial objective optimizes the wrong quantity, a hidden constraint makes the requested solution infeasible, or the verifier encodes an invalid specification. The central design question is therefore not only how an agent reaches a goal, but how the runtime determines whether to persist, replan, or ask the user to clarify which goal and constraints should govern the next round.

Argus addresses this question by coupling persistent self-evolution with independent verification. Persistence alone can compound a false assumption. Verification alone can repeatedly diagnose the same issue without improving later work. Together,

they create a reviewed frontier that supports an explicit choice: persist when evidence still supports the contract, pivot the technical approach when it does not, or surface evidence and alternatives when the correct objective or constraint requires user judgment.

This is not goal drift. Equation 3 requires evidence, intent preservation, authority, and provenance. The standing intent can remain stable while the operational objective changes substantially. A user asking for a particular optimization may learn that latency is dominated elsewhere. A verification engineer may learn that the implementation satisfies the written spec but the spec violates the intended safety property. A mathematician may replace an overbroad conjectural route with a sequence of reachable frontier claims. In each case, fidelity to the initial wording would be less aligned than a verified, authorized refinement.

7.2 System-Level Performance

The seven-benchmark suite shows one runtime operating across software repair, GPU-kernel optimization, model training, research tasks, data synthesis, and mathematical proof. On SWE-Bench Pro, Argus reaches approximately 78% accuracy compared with 59% for Direct Copilot at $1.41 \times$ aggregate Tokens. The comparison pairs additional computation for planning, execution, and review with a 19-point accuracy difference.

The longitudinal run shows how this cost changes with accumulated state. Mature W19–22 uses 21% fewer solve input Tokens and 15% less active time per task than startup W1–6. The curve remains non-monotone: W13–18 combines low Token use with longer execution, while W23–24 contains a late concentration of difficult tasks. The operating profile changes over time as shared state expands, while task-dependent variation remains visible.

7.3 Autonomous Persistence with Selective User Involvement

User involvement is selective rather than continuous. The evidence shows that the same harness remains productive when the user supplies intent and constraints but does not occupy the inner loop. The six paper campaigns span 639.55 aggregate campaign-hours, 254 missions, 576 Engineer rounds, and 89 session rolls, with all six canonical pipelines reaching submission completion. The FLA case records an operator-supplied target and hardware constraint followed by autonomous profiling, implementation, benchmarking, and correctness validation; external human review then supplied a numerical-stability correction before merge. The 731-task SWE-Bench Pro trajectory provides a third, repeated-task view.

These observations do not prove that human input is unnecessary. They show a more specific property: once the task contract is sufficiently actionable, continuity, experimentation, verification, recovery, and approach self-evolution are carried by the harness rather than by continual operator prompting. Human expertise can enter where it has the highest marginal value—contract ambiguity, domain judgment, and material tradeoffs—without becoming a hidden scheduler for ordinary rounds.

7.4 Honest Stopping as a Reliability Property

Informal trials with internal domain experts repeatedly highlighted a qualitative property that endpoint accuracy does not capture: Argus often stops at an appropriate boundary instead of manufacturing progress. This feedback is anecdotal rather than a controlled user study, but it identifies an observable runtime property supported by both code and retained trajectories.

The runtime separates `done`, `continue`, `blocked`, `replan_requested`, `no_progress`, and `infra_blocked`; it also keeps execution status, review status, Stage certification, interruption kind, and resumability as independent dimensions. Unknown progress is not guessed from prose: only the Reviewer’s explicit structured negative judgment advances the semantic stall counter. In the archived SWE-Bench Pro run, 35 tasks were blocked on first external review rather than accepted. Across the six paper campaigns, Reviewer issued 88 blocked verdicts and Manager performed 16 Stage rollbacks, yet all six pipelines eventually completed. The mathematical trace retains a falsified route as a useful result rather than relabeling it as theorem progress.

Honest stopping is therefore not conservative refusal. The same evidence contains 43 SWE-Bench revision requests, 34 later official-verifier successes, 22 strict rescues, and repeated paper-stage recovery. The harness stops the current claim or approach while preserving enough state to revise, replan, or resume. This balance—neither endless continuation nor premature success—is central to unattended long-horizon reasoning.

7.5 Reviewer as Selective Error Correction and Contract Critic

The Reviewer requests revision on 43 SWE-Bench Pro tasks. After another Engineer round, 34 pass the official verifier and 22 complete the stricter `continue`→`revision`→`done` loop. These trajectories show the Reviewer functioning as an error-correction stage rather than a terminal commentary layer.

At longer horizons, error correction also applies to the task contract. A failed check can mean that the artifact is wrong, but it can also reveal an inconsistent constraint, an underpowered evaluator, or an acceptance criterion disconnected from expert value. Reviewer does not have unilateral authority to invent the replacement goal. Its role is to ground the contradiction, reject unsupported completion, and route the resulting decision to Planner, Manager, or the user at the proper boundary.

Reviewer routing is adaptive: Reviewer-routed tasks use more Tokens and active time than historically self-reviewed tasks. The observed routing concentrated independent review on the more resource-intensive part of that archived workload; it should

not be read as the current routing policy, which independently reviews every normal round.

7.6 Artifact-Level Outcomes

The merged FLA kernel and ACE-1 prototype illustrate two different forms of real-world output. FLA PR #1045 crossed an external repository boundary, received human review, was revised for numerical stability, and landed upstream. ACE-1 crossed a hardware-artifact boundary by producing synthesizable RTL and a gate-level netlist, but the supplied report material exposes only the project-level score and artifact description. The first is externally auditable adoption; the second is a promising operator-reported system artifact pending reproduction from raw design and scoring inputs.

7.7 Research Production as Contract Refinement

The six-paper case study extends the endpoint results into a complete research lifecycle. All six canonical pipelines reach submission completion across six domains and two venue formats, but none follows a one-pass path. Reviewer revision verdicts, Stage rollbacks, and session rolls are common rather than exceptional. The final manuscripts therefore demonstrate persistence across research framing, experimentation, writing, review, and packaging—not merely the ability to emit paper-like prose.

The multimodal-hallucination trace makes the control mechanism explicit. Reviewer gates prevent seven weak method routes from being scaled into an unsupported positive claim; Planner reframes the accumulated failures as a diagnostic study; Engineer then completes the 15-cell canonical matrix; and Manager accepts two late rollbacks when submission checks expose scorer-provenance and telemetry defects. The role separation therefore changes the scientific trajectory: failure is retained as information, while only reviewed state controls the next Stage.

The reframing is the empirical analogue of $K_t \rightarrow K_{t+1}$. The broad intent—produce a valid contribution about mitigation behavior—survives, while the operational objective changes from demonstrating a successful method to characterizing why candidate wrappers fail. Because the no-go measurements remain linked to the new claim, this transition is auditable. The system does not reward it merely for changing direction; it rewards the verified information gain and the more defensible contract that results.

The case also separates artifact completion from process certification. Five final assurance snapshots pass; the compositional-matching project retains a stale blocked assurance snapshot after its canonical pipeline and final PDF complete. This is a useful systems distinction: derived quality metadata must be reconciled with the authoritative Stage state before it can serve as a public certificate.

7.8 Persistent State Makes Pivots Cumulative

The long-lived object in Argus is the campaign state rather than a provider transcript. Accepted results, failed routes, tools, skills, and current blockers remain available after session and process restarts. Equation 13 formalizes the information advantage of this record, while Equation 15 explains why the runtime keeps both a complete event history and a bounded working checkpoint.

This mechanism gives concrete meaning to *fixed-model runtime self-evolution*. The base model remains unchanged, but an accepted mission can change the starting point, available procedures, verification obligations, or routing policy of later work. Rejected branches also participate when they are retained as evidence-backed exclusions. Evolution therefore occurs in the runtime’s reviewed state and control policy, not in model weights or an unconstrained summary of prior conversation.

Agent-side evolution improves execution under the current contract, but it cannot decide every contract change. The state in Equation 9 also includes X_t and K_t . A stronger Skill may reduce future execution cost; a stronger verifier may instead show the user that the current target is invalid; a negative result may do both. Retaining that evidence makes a pivot cumulative: later work starts from a better specification of what should be attempted next and why, rather than discarding the failed route or silently inventing a new goal.

The two longitudinal views expose complementary parts of this claim. The SWE-Bench Pro trajectory measures how solve Tokens and active time vary as shared state accumulates, while the mathematical campaign shows direct semantic reuse: later missions consume earlier definitions, bounds, rejected routes, and bridge lemmas. Neither trace isolates a single update mechanism. A matched experiment that replays the same tasks with frozen versus accumulated runtime state would be required to attribute a causal gain to memory, Skills, verification, or routing separately.

The current traces also do not directly measure how much a user’s understanding improved. They expose the evidence, contract changes, and decision points that make such improvement possible. A prospective study should measure expert calibration, time to identify the true constraint, contract revision quality, and final task value against fixed-prompt and unverified-agent baselines.

7.9 Implications for Engineering and System Verification

Engineering requests frequently specify a desired implementation before the user knows the binding constraint. A verification-centered loop can treat prototypes and benchmarks as instruments for eliciting the real objective. The useful intermediate result may be a bottleneck decomposition, a rejected architecture, or a revised latency–memory tradeoff rather than the initially requested patch. Persistent state ensures that these discoveries constrain later implementation instead of being forgotten between

sessions.

System verification makes the need even sharper because there are at least three competing hypotheses after a failed obligation: the code is wrong, the specification is wrong, or their mapping is wrong. A completion-oriented agent is biased toward editing code until the current checker passes. Argus’s intended workflow retains the counterexample, audits both sides of the relation, strengthens the verifier when needed, and escalates normative ambiguity to the user. Passing a checker is evidence only under a contract whose assumptions and safety meaning have also survived review.

7.10 From Runtime Learning to Model Training

The same trajectories can be transformed into supervised examples, preference pairs, and reinforcement-learning transitions. AgentInstruct, Agent-FLAN, and Agent Lightning provide complementary mechanisms for this conversion (Mitra et al., 2024; Chen et al., 2024; Luo et al., 2025). Argus contributes a structured source stream linking objectives, tool actions, measurements, revisions, and retained-state updates. A future training stage can use this stream to internalize recurring planning and verification patterns while the runtime continues to provide long-horizon coordination.

7.11 Vertical Research in Mathematics

The Erdős–Gyárfás campaign exposes the four-role organization at research depth. The Manager preserves the objective and stage, the Planner converts the current frontier into bounded theorem tasks, the Engineer produces proofs and computational checks, and the Reviewer controls admission into durable state.

The campaign retains one falsified route and six accepted theorem-frontier updates. Later missions improve a bound and close two graph-realizability gaps by starting from the reviewed ledger rather than restarting from the original conjecture. This is the same runtime mechanism seen in software repair, applied to a domain where progress is expressed through definitions, lemmas, counterexamples, and proofs.

7.12 Connecting Theory and Results

Table 6 connects the verified persistence-and-pivot model to the measurements reported in Sections 6 and 6.5.

Theory component	Measurement	Observed result	Interpretation
Process-data dominance	Mathematical trajectory	One dead route and six dependent theorem updates	The retained process changes the starting state of later missions.
Intermediate value R_t	Mathematical frontier	One falsification, one stronger bound, and bridge results before terminal closure	Reviewed intermediate states carry task-native value and redirect later missions.
Review gate	43 Reviewer revision requests	34 verifier recoveries; 22 strict rescues	Independent review redirects initially unacceptable work into another implementation round.
Honest stopping	SWE-Bench and paper trajectories	35 first-review blocks; 88 paper-review blocks; 16 Stage rollbacks	The runtime records non-success without discarding the state needed for recovery.
Reuse value G_L	Startup W1–6 versus mature W19–22	21% fewer solve input Tokens; 15% less active time	Later operation reaches a lower resource regime while shared state expands.
Contract refinement K_t	Paper-production trajectory	Seven early no-go rollbacks and a diagnostic re-framing	The operational claim changes under retained evidence while the standing scientific intent remains stable.
Recovery under persistent state	Paper-production trajectory	Two late submission repairs and 29 session rolls	The campaign repairs evidence without losing accepted results or restarting the manuscript.
Token conversion	Rounds 12–17 role attribution	56.0% reasoning, 55.4% verification, and 55.6% action direct shares	The research trace separates frontier work, coordination, and interrupted work.

Table 6. Connection between verified persistence-and-pivot theory and measured system behavior. Detailed substitutions appear in Appendix D.

8 Limitations and Future Work

The benefit of user-guided pivots is not yet causally evaluated. The retained traces show evidence-backed changes to operational objectives, constraints, and verification obligations, but they do not include a prospective human-subject comparison. We therefore do not claim that the current experiments prove improved expert understanding. A direct evaluation should compare Argus with fixed-prompt agents and ordinary interactive assistants on expert calibration, time to identify the binding constraint, quality of contract revisions, number of repeated misconceptions, and final task value. Likewise, internal-expert observations about appropriate stopping motivate the analysis in Section 7 but are not reported as a satisfaction rate or comparative user-study result.

Contract refinement can still fail. Evidence, provenance, and authority reduce goal drift but do not eliminate it. An expert may approve a poorly framed tradeoff, Reviewer may miss a contradiction, or the recorded standing intent may omit a tacit

requirement. Future work should test adversarial contract changes, invariant-preservation checks, disagreement handling, and rollback of harmful refinements. High-impact domains also require explicit human sign-off policies rather than relying on Manager judgment alone.

The task-contract notation is an analytical abstraction over several persisted runtime surfaces, not a claim that the implementation currently exposes one atomic K_i object or a universal automated user-admission API. A future implementation should unify those fields and make authorization and rollback directly queryable.

Verification is only as sound as its evidence boundary. An executable test, formal checker, benchmark, or model Reviewer can all encode the wrong property. The framework makes verifier revision possible but does not confer correctness automatically. System-verification studies should separately seed bugs in code, bugs in specifications, and mismatches between them, then measure whether the runtime localizes the faulty side instead of merely making the current checker pass.

Attribution and task sequence. The SWE-Bench Pro experiment compares the complete Argus runtime with Direct Copilot. Reviewer routing is adaptive, the 22 Waves follow one task order, and per-Wave Direct-Copilot Token/time records are unavailable. The current results therefore characterize whole-system behavior over this sequence. Matched frozen-state runs, randomized task orders, and randomized review routing will separate the contributions of persistent state, role separation, and review.

Generality across models, domains, and evaluators. The seven benchmark arenas use different metrics, hardware, and execution protocols, while the mathematical analysis follows one vertical campaign. Results are interpreted in their native units rather than combined into a universal score. Repeated campaigns, additional backbones, and external domain review will measure how the runtime transfers across task families and research standards.

Metric calibration and model transfer. The efficiency analysis labels whole role calls by mission purpose; finer segment labels would sharpen the reasoning, action, and verification factors. Reviewer sensitivity and false-acceptance rate require randomized routing with an external verifier, and the counterfactual reuse value G_L requires paired future tasks with and without a state update. The runtime-to-model path is currently a training direction: held-out trajectory splits and scaffold-removal studies will measure how much of the long-horizon control loop can be internalized by the model.

Paper-production trajectories are observational. The six paper projects were selected from one shared research environment and do not include a matched human-authoring baseline. Campaign-hours overlap in calendar time, stored review snapshots are model-generated rather than independent human reviews, and runtime logging evolved across projects. The case study therefore supports end-to-end lifecycle and recovery claims, not paper acceptance, scientific novelty, or superiority to human research teams. Recorded costs in the public data also exclude utility calls without a priced event.

Artifact evidence levels. The FLA case is linked to a public merged pull request and archived timing and correctness evidence. The ACE-1 description and 75/80 score come from the supplied project presentation; the report bundle does not include enough raw PPA and throughput material for independent recomputation. ACE-1 is therefore evidence of an artifact-producing chip-design workflow, not a reproduced SOTA comparison.

9 Conclusion

Argus is a verification-centered framework for persistent and pivot-capable long-horizon reasoning. Its premise is that success does not come from rushing in one direction for as long as possible. The runtime must persist while evidence supports the current contract, but pivot when verification reveals a failed route, hidden constraint, invalid evaluator, or incorrectly framed objective. Treating such evidence as failure to follow the initial prompt is too rigid; allowing an agent to rewrite its goal freely is unsafe. Argus instead makes each material pivot a persistent, evidence-backed, authorized transition.

The framework separates standing intent from the revisable task contract $K_i = (t, o_i, c_i, v_i)$. Manager preserves intent and owns contract and Stage transitions; Planner converts the reviewed frontier into bounded work; Engineer produces artifacts and measurements; and a fresh independent Reviewer controls admission after every normal round. Users enter at consequential decision boundaries, where evidence and alternatives can refine value-laden goals or constraints. Verification is therefore not a final checkbox. It determines whether the runtime should continue, revise the approach, stop, or surface evidence and alternatives for user clarification or authorization.

Verifier-gated runtime self-evolution remains the execution engine of this framework. Across Stages, Π_i records attempted approaches, measurements, accepted and rejected routes, Reviewer decisions, skills, and verifier changes. Later missions retrieve this reviewed state to choose a better approach between the current goal and constraints while model weights remain

fixed. A verified pivot can refine the task contract; self-evolution prevents the runtime from restarting its search whenever that contract or Stage changes.

The empirical results remain strong under this framing. Across seven benchmark arenas, Argus reaches task-native competitive outcomes; on SWE-Bench Pro it attains approximately 78% accuracy versus 59% for Direct Copilot at $1.41\times$ aggregate Tokens. Mature operation uses 21% fewer solve input Tokens and 15% less active time than startup, while Reviewer intervention produces 34 verifier recoveries and 22 strict rescue loops. These endpoint results establish that persistence and verification do not prevent delivery.

The longitudinal cases show why endpoints are not enough. The mathematical campaign retains one falsified route and six proof-backed frontier updates; a counterexample, bridge lemma, or stronger bound has value before terminal proof. The six-paper case survives 16 Stage rollbacks, and its representative campaign turns seven verified no-go routes into a 4,500-row negative-results study rather than manufacturing a positive claim. These are examples of intermediate reward and contract refinement: the standing intent is preserved while the operational target becomes more faithful to evidence.

The same harness does not require step-by-step human supervision once a contract is actionable. Six paper pipelines span 254 missions and 576 Engineer rounds; the FLA case proceeds from an operator-supplied target to an autonomously implemented and validated kernel before external review. Equally important, the runtime can decline to certify progress: structured blocked, stalled, incomplete, and replan outcomes, together with 35 first-review blocks, 88 paper-review blocks, and 16 Stage rollbacks, show honest stopping without preventing later recovery.

Artifact-level outcomes extend the same loop beyond benchmark endpoints. A TileLang RWKV6 kernel was revised under external human review and merged into Flash Linear Attention. The supplied ACE-1 record reports synthesizable Sky130 NPU artifacts and a 75/80 project score; absent raw reproduction inputs, this report treats that result as operator-reported rather than independently validated.

The current study does not yet prove a causal improvement in user understanding; that requires prospective collaboration studies. It does establish the runtime mechanisms, evidence surfaces, and observed trajectories needed to test that claim. More broadly, the retained objectives, contract revisions, actions, measurements, verifier decisions, and state transitions define a natural interface to supervised, preference-based, and reinforcement learning. The long-term goal is not an agent that merely persists longer on the initial prompt. It is a verified system that knows when to keep going, when to pivot, and when the meaning of success must be clarified or authorized by the user.

A Detailed SWE-Bench Pro Results

The main text reports the startup-to-mature comparison. Table 7 shows every analysis window, including the composition shift and late difficult-task period.

Window	Interpretation	Tasks	Input/task	Active min/task	Token index
W1–6	Startup	120	2.95M	8.52	100
W7–12	Early reuse	140	2.07M	7.28	70
W13–18	Composition shift	151	1.47M	10.42	50
W19–22	Mature	158	2.33M	7.25	79
W23–24	Late difficult tasks	49	3.72M	9.01	126

Table 7. Task-weighted Argus window statistics. Token index normalizes W1–6 to 100.

Wave	Tasks	Solve input/task	Agent s/task	Skills	Wiki
1	20	2.839M	545.6	34	27
2	20	3.263M	642.6	54	41
3	20	3.609M	569.6	68	52
4	20	2.908M	490.7	89	64
5	20	2.832M	452.0	106	81
6	20	2.246M	366.8	127	94
7	20	2.564M	519.8	148	101
8	20	2.882M	465.8	168	115
9	20	2.615M	394.2	190	129
10	20	2.204M	377.1	208	143
11	20	2.041M	396.9	225	155
12	40	1.099M	451.5	260	171
14	32	1.327M	527.0	290	191
16	39	1.564M	623.1	305	206
17	40	1.537M	668.5	306	219
18	40	1.432M	662.8	333	237
19	40	2.424M	431.0	365	258
20	40	2.129M	400.5	404	284
21	39	2.165M	427.9	413	305
22	39	2.593M	481.3	434	323
23	38	3.873M	565.6	471	345
24	11	3.186M	455.4	478	352

Table 8. All completed Argus Wave summaries used in the longitudinal analysis. Two incomplete Waves are omitted from grouped means.

B Reviewer Intervention Details

Table 9 separates routing decisions from recovery outcomes. The first three rows partition the independent-Reviewer trajectories; the final two rows measure recovery after a revision request.

Reviewer mechanism event	Tasks	Interpretation
Independent Reviewer invoked	466	63.7% of 731
Historical Engineer self-review route	265	36.3% of 731
Accepted on first external review	388	terminal done
Blocked on first external review	35	terminal blocked
External Reviewer requested revision	43	at least one continue
Official verifier success after revision	34	79.1% of revision requests
Strict review-loop rescue	22	51.2% of revision requests

Table 9. Reviewer routing and recovery in the archived SWE-Bench Pro runtime. Strict rescue requires an external Reviewer **continue** followed by a later Reviewer **done**; current production rounds all use independent review.

C Representative Vertical Trace Details

Table 10 expands the role handoffs summarized in Figure 5. The table deliberately foregrounds Agent authority and durable outputs rather than the mathematical statements.

Phase	Manager / Planner	Engineer	Reviewer	Durable output
Recover and scope	Manager restores the campaign; Planner screens candidate problems.	Grounds sources and creates deterministic checks.	Verifies problem fidelity and accepts the scoped objective.	Persistent objective, selected problem, initial claim record.
Test a route	Planner chooses the cheapest decisive falsification mission.	Runs the witness verifier and records the failed reduction.	Accepts the negative result and fixes its scope.	Certified dead branch and revised next action.
Produce proof package	Planner requires theorem, proof, and verification artifacts.	Writes proof, validation record, lemma graph, and checkpoint.	Returns continue when checks are incomplete, then rechecks.	Reviewer-accepted bounded result package.
Strengthen	Planner requires later work to consume and advance retained state.	Improves one result and proves bridge/obstruction lemmas.	Checks correctness, strict progress, and claim boundaries.	Updated accepted state plus named remaining blockers.
Continue	Manager advances or rolls back from the verdict; Planner selects the next blocker.	Starts the next mission from retained artifacts.	Updates the frontier status and next blockers.	Inspectable research trace and reusable cross-session state.

Table 10. Role-resolved phases in the reported Erdős–Gyárfás campaign. The table follows how each role changes the shared research state.

The supplementary trace bundle contains the role/claim table, calculation map, editable figure source, and provenance. The source campaign retains the complete artifacts and event history.

Representative proof artifact (abridged)

For fixed roots x, y , let $\mathcal{P}_4(x, y)$ be the length-4 paths from x to y , and let $\mathcal{F}$ be their internal three-vertex sets. In a graph with no C_8 , any two members of $\mathcal{F}$ intersect: two paths with disjoint internal vertices would join to form an 8-cycle. A separate no- C_4 case analysis shows that any fixed pair of internal vertices occurs in at most three path sets. Assume that the path family has no common internal blocker. Choose $E = \{a, b, c\} \in \mathcal{F}$. For each $t \in E$, some member E_t omits t . Every set containing t must meet the three vertices of E_t ; the pair-occurrence bound therefore gives $d(t) \leq 9$. Hence

$$d(a) + d(b) + d(c) \leq 27.$$

Every member of $\mathcal{F}$ meets E , while E itself contributes three incidences rather than one, so

$$|\mathcal{F}| + 2 \leq d(a) + d(b) + d(c) \leq 27.$$

Thus $|\mathcal{P}_4(x, y)| = |\mathcal{F}| \leq 25$: the paths share a common internal blocker or there are at most 25 of them. This excerpt illustrates the kind of proof artifact produced and reviewed by the workflow; the complete case analysis and computational checks remain in the source bundle.

D Process-Theory Calculations

Table 11 records the complete substitutions behind the process quantities summarized in the main text. The mathematical window contains 332,274 reasoning-output Tokens over 9,620.5 active seconds. The strict and inclusive density values use the two attribution policies defined in Equation 8. The process-compression row is a separate campaign snapshot.

Theory quantity	Data substitution	Observed value	Interpretation
Strict $\hat{p}_l$	$332,274/9,620.5 \times 0.5598 \times 0.5556 \times 0.5538$	5.95 effective Token-eq./s	Direct reasoning, action, and verification only.
Inclusive $\hat{p}_l$	$332,274/9,620.5 \times 0.9512 \times 0.9444 \times 0.9043$	28.06 effective Token-eq./s	Includes successful coordination and state-maintenance work.
Reviewer correction	34/43 verifier recoveries; 22/43 strict rescues	79.1%; 51.2%	Post-revision outcomes on SWE-Bench Pro.
Reuse quantity G_L	(2.95 – 2.33)M Tokens/task; (8.52 – 7.25) min/task	0.62M Tokens and 1.27 min per task	Startup–mature difference in the longitudinal run.
Immediate $\widehat{\mathcal{P}}_{\text{PRI}}$	$\lambda_g = 0$; $6/332,274 \times 10^6$	18.1 accepted deltas/M reasoning Tokens	Six accepted theorem-frontier updates in Rounds 12–17.
Process compression	384,370,463/3,240 bytes	118,633× event-history/checkpoint ratio	Event history relative to the active working checkpoint.

Table 11. Detailed calculation of the process-theory quantities. Values retain their task-specific units; the main text reports the corresponding system-level patterns.

E Upstream Kernel Adoption

The kernel-optimization work produced an externally reviewed upstream contribution to Flash Linear Attention. In [FLA PR #1045](#), Argus implemented and optimized an opt-in TileLang RWKV6 dense-bf16, D=64 forward-intra kernel. The submitted H100 evidence reports a forward latency reduction from 0.199 ms to 0.168 ms ($1.18\times$) and a forward-plus-backward reduction from 0.900 ms to 0.747 ms ($1.21\times$), with 13 correctness-gate passes and 14 repository tests.

An FLA collaborator affiliated with Moonshot AI reviewed the generated CUDA, identified a long-sequence numerical-stability issue, and requested block-local exponent centering. After Argus applied the fix and reran the verification suite, the contribution was merged into `fla-org:main` on July 20, 2026 as [commit c70f11c](#). This is evidence of human-reviewed upstream adoption, not external scientific validation of the broader runtime.

F Notation

Symbol	Meaning
o, y_0, f, B	Objective, initial artifact, evaluator, and resource budget
τ_t	Mission trajectory $(s_t, a_t, y_t, e_t, r_t)$
H_t	Persistent runtime state
E_t	Results admitted after mission t
θ_t	Underlying model parameters, fixed in this report
C_t	Effective retained capability set
$\phi(\tau_t)$	Role-resolved mission trace $(m_t, p_t, x_t, r_t, \Delta H_t)$
$\rho_l(T)$	Dense-intelligence density over horizon T
η_r, η_a, η_v	Measured reasoning, action, and verification efficiency factors
$\mathcal{R}_q(X)$	Minimum downstream decision risk when observing X
ψ_b^*	Decision-useful process compression under context budget b
$G_L(\Delta H_t)$	Discounted counterfactual reuse value over L future tasks
$\mathcal{P}_{\text{PRI}}(W)$	Verified reusable process-intelligence yield per Token
α, β	Reviewer sensitivity and false-acceptance rate
$V_R(T)$	Conceptual accumulated research value
R_{tok}	Aggregate Argus-to-Direct-Copilot Token ratio
$\bar{t}_w^{\text{solve}}$	Mean solve input tokens in Wave w
$\bar{T}_w^{\text{agent}}$	Mean active workflow time in Wave w

Table 12. Notation used in the report.

Machine-readable aggregates and editable figure sources are released with the report. They are kept outside the narrative so that implementation metadata does not obscure the scientific results.

References

- Argus Team. Argus: A self-evolving research agent. Project website, 2026. URL <https://argusbot.cn/>.
- Jinheon Baek, Sujay Kumar Jauhar, Silviu Cucerzan, and Sung Ju Hwang. ResearchAgent: Iterative research idea generation over scientific literature with large language models. *arXiv preprint arXiv:2404.07738*, 2024. URL <https://arxiv.org/abs/2404.07738>.
- David Blackwell. Equivalent comparisons of experiments. *The Annals of Mathematical Statistics*, 24(2):265–272, 1953. doi: 10.1214/aoms/1177729032. URL <https://doi.org/10.1214/aoms/1177729032>.
- Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, et al. MLE-bench: Evaluating machine learning agents on machine learning engineering. *arXiv preprint arXiv:2410.07095*, 2024. URL <https://arxiv.org/abs/2410.07095>.
- Zehui Chen, Kuikun Liu, Qiuchen Wang, Wenwei Zhang, Jiangning Liu, Dahua Lin, Kai Chen, and Feng Zhao. Agent-FLAN: Designing data and methods of effective agent tuning for large language models. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 9354–9366. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.findings-acl.557. URL <https://aclanthology.org/2024.findings-acl.557/>.
- Xiang Deng, Jeff Da, Edwin Pan, Yannis Yiming He, Charles Ide, Kanak Garg, Niklas Lauffer, Andrew Park, Nitin Pasari, Chetan Rane, Karmini Sampath, Maya Krishnan, Srivatsa Kundurthy, Sean Hendryx, Zifan Wang, Vijay Bharadwaj, Jeff Holm, Raja Aluri, Chen Bo Calvin Zhang, Noah Jacobson, Bing Liu, and Brad Kenstler. SWE-Bench Pro: Can AI agents solve long-horizon software engineering tasks? *arXiv preprint arXiv:2509.16941*, 2025. URL <https://arxiv.org/abs/2509.16941>.
- Alireza Ghafarollahi and Markus J. Buehler. SciAgents: Automating scientific discovery through multi-agent intelligent graph reasoning. *arXiv preprint arXiv:2409.05556*, 2024. URL <https://arxiv.org/abs/2409.05556>.
- Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, et al. Accelerating scientific discovery with co-scientist. *arXiv preprint arXiv:2502.18864*, 2025. URL <https://arxiv.org/abs/2502.18864>.
- Zhengyao Jiang, Dominik Schmidt, Dhruv Srikanth, et al. AIDE: AI-driven exploration in the space of code. *arXiv preprint arXiv:2502.13138*, 2025. URL <https://arxiv.org/abs/2502.13138>.
- Carlos E. Jimenez, John Yang, Alexander Wettig, et al. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2024. URL <https://arxiv.org/abs/2310.06770>.
- Jiajie Jin, Yuyang Hu, Kai Qiu, Qi Dai, Chong Luo, et al. Toward generalist autonomous research via hypothesis-tree refinement. *arXiv preprint arXiv:2606.11926*, 2026. URL <https://arxiv.org/abs/2606.11926>.
- Keller Jordan and contributors. Modded-NanoGPT: The nanogpt speedrun. GitHub repository, 2024. URL <https://github.com/KellerJordan/modded-nanogpt>.
- Andrej Karpathy. nanochat: The best ChatGPT that \$100 can buy. GitHub repository, 2025. URL <https://github.com/karpathy/nanochat>.
- Andrej Karpathy. autoresearch: AI agents running research on single-GPU nanochat training automatically. GitHub repository, 2026. URL <https://github.com/karpathy/autoresearch>.
- Thomas Kwa, Ben West, Joel Becker, et al. Measuring AI ability to complete long software tasks. *arXiv preprint arXiv:2503.14499*, 2025. URL <https://arxiv.org/abs/2503.14499>.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023. URL <https://arxiv.org/abs/2305.20050>.
- Edward Lin, Sahil Modi, Siva Kumar Sastry Hari, et al. SOL-ExecBench: Speed-of-light benchmarking for real-world GPU kernels against hardware limits. *arXiv preprint arXiv:2603.19173*, 2026. URL <https://arxiv.org/abs/2603.19173>.
- Xiao Liu, Hao Yu, Hanchen Zhang, et al. AgentBench: Evaluating LLMs as agents. *arXiv preprint arXiv:2308.03688*, 2023. URL <https://arxiv.org/abs/2308.03688>.
- Chris Lu et al. The AI scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292*, 2024. URL <https://arxiv.org/abs/2408.06292>.

- Xufang Luo, Yuge Zhang, Zhiyuan He, Zilong Wang, Siyun Zhao, Dongsheng Li, Luna K. Qiu, and Yuqing Yang. Agent lightning: Train any AI agents with reinforcement learning. *arXiv preprint arXiv:2508.03680*, 2025. URL <https://arxiv.org/abs/2508.03680>.
- Alisia Lupidi, Bhavul Gauri, Thomas Simon Foster, et al. AIRS-Bench: A suite of tasks for frontier AI research science agents. *arXiv preprint arXiv:2602.06855*, 2026. URL <https://arxiv.org/abs/2602.06855>.
- Chang Ma, Junlei Zhang, Zhihao Zhu, et al. AgentBoard: An analytical evaluation board of multi-turn LLM agents. *arXiv preprint arXiv:2401.13178*, 2024. URL <https://arxiv.org/abs/2401.13178>.
- Arindam Mitra, Luciano Del Corro, Guoqing Zheng, Shweti Mahajan, Dany Rouhana, Andres Cudas, et al. AgentInstruct: Toward generative teaching with agentic flows. *arXiv preprint arXiv:2407.03502*, 2024. URL <https://arxiv.org/abs/2407.03502>.
- Alexander Novikov, Ng  n V  , Marvin Eisenberger, et al. AlphaEvolve: A coding agent for scientific and algorithmic discovery. *arXiv preprint arXiv:2506.13131*, 2025. URL <https://arxiv.org/abs/2506.13131>.
- Charles Packer, Sarah Wooders, Kevin Lin, et al. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023. URL <https://arxiv.org/abs/2310.08560>.
- Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, et al. Generative agents: Interactive simulacra of human behavior. *arXiv preprint arXiv:2304.03442*, 2023. URL <https://arxiv.org/abs/2304.03442>.
- Weitong Qian, Beicheng Xu, Zhongao Xie, et al. AutoSci: A memory-centric agentic system for the full scientific research lifecycle. *arXiv preprint arXiv:2605.31468*, 2026. URL <https://arxiv.org/abs/2605.31468>.
- Bernardino Romera-Paredes, Mohammadamin Barekatain, Alexander Novikov, et al. Mathematical discoveries from program search with large language models. *Nature*, 625:468–475, 2024. doi: 10.1038/s41586-023-06924-6. URL <https://doi.org/10.1038/s41586-023-06924-6>.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dess  , et al. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023. URL <https://arxiv.org/abs/2302.04761>.
- Samuel Schmidgall and Michael Moor. AgentRxiv: Towards collaborative autonomous research. *arXiv preprint arXiv:2503.18102*, 2025. URL <https://arxiv.org/abs/2503.18102>.
- Samuel Schmidgall, Yusheng Su, Ze Wang, et al. Agent laboratory: Using LLM agents as research assistants. *arXiv preprint arXiv:2501.04227*, 2025. URL <https://arxiv.org/abs/2501.04227>.
- Noah Shinn et al. Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*, 2023. URL <https://arxiv.org/abs/2303.11366>.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*, 2024. URL <https://arxiv.org/abs/2408.03314>.
- Qiong Tang, Tianxiang Sun, Xiangkun Hu, et al. FARS: A fully automated research system deployed at scale. *arXiv preprint arXiv:2606.31651*, 2026. URL <https://arxiv.org/abs/2606.31651>.
- Guanzhi Wang et al. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023. URL <https://arxiv.org/abs/2305.16291>.
- Jiayu Wang, Weijiang Lv, Bowen Fu, et al. Act as a real researcher: A suite of benchmarks evaluating frontier LLMs and agentic harnesses in research lifecycle. *arXiv preprint arXiv:2606.07462*, 2026. URL <https://arxiv.org/abs/2606.07462>.
- Xingyao Wang, Boxuan Li, Yufan Song, et al. OpenHands: An open platform for AI software developers as generalist agents. *arXiv preprint arXiv:2407.16741*, 2024a. URL <https://arxiv.org/abs/2407.16741>.
- Zora Zhiruo Wang, Jiayuan Mao, Daniel Fried, and Graham Neubig. Agent workflow memory. *arXiv preprint arXiv:2409.07429*, 2024b. URL <https://arxiv.org/abs/2409.07429>.
- Yixuan Weng, Minjun Zhu, Guangsheng Bao, Hongbo Zhang, Jindong Wang, Yue Zhang, and Linyi Yang. CycleResearcher: Improving automated research via automated review. In *International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=bjcsVLoHYs>.

- Hjalmar Wijk, Tao Lin, Joel Becker, et al. RE-Bench: Evaluating frontier AI R&D capabilities of language model agents against human experts. *arXiv preprint arXiv:2411.15114*, 2024. URL <https://arxiv.org/abs/2411.15114>.
- Wujiang Xu, Zujie Liang, Kai Mei, et al. A-MEM: Agentic memory for LLM agents. *arXiv preprint arXiv:2502.12110*, 2025. URL <https://arxiv.org/abs/2502.12110>.
- Yutaro Yamada, Robert Tjarko Lange, Cong Lu, et al. The AI scientist-v2: Workshop-level automated scientific discovery via agentic tree search. *arXiv preprint arXiv:2504.08066*, 2025. URL <https://arxiv.org/abs/2504.08066>.
- John Yang et al. SWE-agent: Agent-computer interfaces enable automated software engineering. *arXiv preprint arXiv:2405.15793*, 2024. URL <https://arxiv.org/abs/2405.15793>.
- Shunyu Yao et al. ReAct: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2023. URL <https://arxiv.org/abs/2210.03629>.