


Autonomous Research Generation and Understanding System

Dense Intelligence for an Expanding Research Frontier

ARGUS - TECHNICAL SPINE

Every run expands the frontier.



Argus is a research infrastructure that operates a capable coding-and-reasoning model as an autonomous research system over long horizons — hours to days — against real, public benchmarks and open research problems. This report documents the system as engineered: a three-plane architecture, the research runtime and mission lifecycle, the state and knowledge systems, the agent decision interfaces, the experiment and resource infrastructure, the reliability and recovery model, and the evidence architecture that keeps every measurement auditable. It reports the current public results across six arenas and a 41-paper research portfolio, each a scoped claim under its stated protocol and evidence tier. It is an engineering report, not a results paper: it does not claim universal state of the art, an accepted autonomous publication, or correctness outside the stated protocols.

Contents

1	Executive Thesis	3
2	Dense Intelligence	5
3	Why Episodic Agents Stall	7
4	Argus Life and the Technical Spine	9
5	Four Roles, Three Planes, One Research Runtime	10
6	Mission Lifecycle and Durable State	14
7	Evidence and Independent Review	17
8	Reliability, Resources, and Operation	18
9	Public Results Across Six Arenas	21
10	Research Portfolio: 41 Papers Across Six Programs	23
11	Limitations and Open Engineering Problems	24
12	Roadmap	26
A	Key Interfaces and Status Taxonomy	27
B	Persisted State Map	27
C	Event Taxonomy	28
D	Evidence Map	28
E	Paper-Inventory Summary	29

ACT I · THE DENSE-INTELLIGENCE PROBLEM

Research intelligence is sparse in time.

Long-horizon research requires continuity, execution, verification, and durable state—not merely a longer context window.

1 Executive Thesis

Argus sustains dense research intelligence, turns each run’s evidence into runtime capability, and carries that capability into the next unknown task. Every run expands the frontier.

Model intelligence is sparse in time. A capable coding-and-reasoning model is brilliant for the length of a single call and then stops: the reasoning that produced an insight evaporates when the context is dropped or lossily compacted, and the next call begins again with no durable trace of what was learned. Long-horizon research is the opposite of a single call. It is thousands of coupled decisions sustained over hours to days, where the value is not in any one clever step but in keeping judgment, execution, and verification connected across all of them. A longer context window does not close this gap; it only postpones the moment the episode ends.

Argus makes intelligence *dense* by running four persistent, model-driven roles as a continuous loop over persisted project state. A Manager fixes intent and lifetime, a Planner decomposes and schedules, an Engineer retrieves and builds and experiments, and a Reviewer inspects the evidence and decides. Because the loop never dissolves back into a single stateless call, decision, execution, and verification stay coupled: the system can carry a thread of research reasoning far past the horizon at which an episodic agent would have forgotten why it started.

Continuity alone is not enough — it must compound. Every run deposits verified evidence, and that evidence updates the system’s *runtime state*: its memory, skills, tools, verifiers, routing, and evaluations. This update is gated by the Reviewer, so only checked results change the runtime, and it happens with the model’s parameters held fixed — $H(t+1) = U(H(t), \text{trajectory}, \text{evidence})$. The system grows more capable not by retraining a model but by accumulating audited, reusable capability around it.

The consequence is out-of-distribution reach. When the next task arrives, the runtime it inherits already carries the skills, tools, and verifiers earned on prior problems, so the next unknown objective does not start from zero. Each expansion of verified capability enlarges the frontier of problems the system can attempt — the expanding-frontier chain of Figure 1.

ARGUS · TECHNICAL SPINE

Every run expands the frontier.




Figure 1. The expanding-frontier spine. An unknown, out-of-distribution objective enters a dense-intelligence runtime driven by four persistent roles; verified work passes an evidence gate; the gate updates durable runtime state; and the enlarged runtime meets the next unknown task from a higher floor. Model parameters remain fixed, and capability is not guaranteed to grow on every run.

What this report claims

Argus is engineered as a persistent, role-separated runtime that keeps decision, execution, and verification coupled over long horizons.

Each run’s reviewer-verified evidence updates durable runtime state — memory, skills, tools, verifiers, routing, evaluations — with model parameters fixed.

Six public arena results, each a scoped claim under its stated protocol, hardware, and evidence tier.

A 41-paper portfolio compared only to human-authored literature, reported as a de-duplicated inventory.

Reliability and auditability by construction: bounded budgets, one completion authority, an append-only evidence tape.

What it does not claim

That the infrastructure is more capable than the model, or substitutes heuristics for the agent’s scientific judgment.

Monotone empirical capability growth: a run may fail, return only a negative result, or add nothing verified.

Universal state of the art, or superiority over all human researchers or other systems.

An accepted autonomous publication, or a count of accepted papers.

Correctness, safety, or performance outside the exact protocols each result is measured under.

Table 1. The claim boundary the rest of the report is held to.

Argus publishes a live public record [1]. This report reproduces its six current arena results (Section 9, Figure 7), each with the protocol and hardware it was measured under and its published reference: NVIDIA SOL-ExecBench (B200, 101 kernels; global rank #6 with two head-to-head wins over Recursive [7]), nanochat on B200 (0.9636 BPB vs. human SOTA 0.9646) and H100 (0.9855 vs. 0.9879 BPB), the nanoGPT speedrun (79.77s vs. the same-device human record 80.18s, $N=10$), AARRI-Bench (63/82, 76.8% vs. the paper-reported best 68.3%), and Arbor’s site-reported system comparison; two of the six — the nanoGPT speedrun and nanochat on B200 — additionally

carry committed artifact-digest records, while the other four remain website snapshots, and that distinction is kept visible throughout (Section 7). Separately, the public research page [2] lists a portfolio of **41 papers** — 35 manuscripts and 6 drafts across six research programs (Section 10, Figure 8) — compared only to human-authored literature and baselines, and reported as a de-duplicated inventory rather than a count of accepted papers.

Design objective, not a measured guarantee

“Every run leaves the system more capable” is the design objective: a run may fail, produce only a negative result, or add no verified capability. The report does not claim monotone empirical capability growth.

2 Dense Intelligence

Design objective

Keep research intelligence *dense* in time: hold decision, execution, and verification coupled and continuous over persisted state, so that a long-horizon program behaves as one sustained line of reasoning rather than a sequence of disconnected bursts.

2.1 Definition

By *dense intelligence* we mean sustained, coupled research work: at every point in a long run the system is making decisions, carrying them into real execution, and verifying the outcome against evidence — and the thread that connects these is never broken by the end of a model call. Sparse intelligence is the default failure mode of a capable model left to run long: it is brilliant in bursts and empty in between, because each burst ends and its working state is discarded. Density is not about being busy; it is about the fraction of elapsed research time in which genuine decisions are being made, executed, and checked, rather than lost to context recovery, idle polling, or unverifiable self-report. Figure 2 contrasts episodic research, where useful work is separated by context-recovery gaps, with the continuous role loop that runs over persisted project state.

DENSE INTELLIGENCE

Continuity is useful only when decision, execution, and verification remain coupled.

Episodic research

useful work separated by context recovery



Argus Life

continuous role loop over persisted project state



conceptual model · not a reported benchmark

Figure 2. Dense intelligence as continuity of coupled work. Episodic research recovers decision and verification repeatedly across context gaps; Argus keeps decision, execution, verification, and state retention coupled over a persisted project. The schematic is a conceptual model, not a reported benchmark.

2.2 Continuity is not token volume

The tempting shortcut is to equate continuity with a larger context window: keep more tokens live and the run will stay coherent. It will not. A longer window lengthens a single episode but does nothing to couple decision to execution to verification, and it still terminates — at which point the accumulated working state is compacted or dropped and the next episode restarts cold. Token volume measures how much text a model can attend to at once; density measures how much of a research horizon is spent on verified, decision-bearing work. A run can consume enormous context and remain sparse — re-reading the same files, re-deriving the same dead ends — because none of that volume is gated by an independent check or retained as reusable capability. Continuity that matters is structural: it comes from persisting state outside the model and cycling roles over it, not from holding more tokens inside the model.

2.3 A conceptual density model

To make the idea precise enough to reason about — and no more — we summarize it as a time-averaged density over a research horizon T :

$$\rho_{\text{DI}}(T) = \frac{1}{T} \int_0^T \lambda(t) \eta_d(t) \eta_x(t) \eta_v(t) dt. \quad (1)$$

The integrand multiplies a rate of research activity by three efficiencies, so that time spent busy but undirected, or active but unverified, contributes little.

Symbol	Meaning
T	The research horizon over which density is averaged — hours to days, not a single call.
$\lambda(t)$	Instantaneous rate of research activity: how much work the system is attempting at time t .
$\eta_d(t)$	Decision efficiency: the fraction of that activity that is a genuine decision advancing the frontier rather than repetition.
$\eta_x(t)$	Execution efficiency: the fraction of decisions carried into real execution on real tools, data, and hardware.
$\eta_v(t)$	Verification efficiency: the fraction of executed work that is independently checked against evidence.

Table 2. The five symbols of the conceptual density model.

Read this way, the episodic failure mode is exactly a collapse of one or more efficiencies to near zero for long stretches: high λ but low η_d (busy, not deciding), or non-trivial η_d and η_x but low η_v (acting, not verifying). The design of Argus is an attempt to hold all three efficiencies away from zero across the whole horizon by construction — persistent roles for decision and verification, real execution for η_x , and an evidence gate for η_v .

2.4 What is and is not measured

The model above is a lens, not an instrument. ρ_{DI} is an explanatory construct, not a reported benchmark metric. It does not, in itself, establish that Argus is universally superior to human researchers or other systems. No number in this report is a measurement of ρ_{DI} : the integrand’s efficiencies are not instrumented as a live scalar, and we make no attempt to publish one. What *is* measured is reported elsewhere under explicit protocols — the six public arena results and the 41-paper portfolio of Section 9 and Section 10, each a scoped claim under its stated evidence tier. Dense intelligence is the design goal those results are evidence *for*; it is not itself a score, and it is not a claim of general superiority.

3 Why Episodic Agents Stall

Design objective

Diagnose why a capable model, run long, degrades into busy sparsity — and derive the design requirements that a system must meet to make research *compound* across runs instead of restarting from zero.

3.1 Model capability is not system capability

It is easy to assume that a more capable model is, by itself, a more capable research system. A stronger model does raise the ceiling of any single step: it reads code more accurately, writes better prose, and reasons further in one call. But a research *system* is judged over thousands of steps sustained across hours to days, and that horizon is governed by properties the weights do not supply: whether the thread of work survives the end of a call, whether claims are checked before they are believed, whether what was learned is retained and reused. Model capability is per-call; system capability is cross-call. Treating them as the same thing is why “give a big model a hard problem and wait” fails — the missing capability lives in the substrate around the model, not in the model.

3.2 Episodic work pays context-recovery cost

Left to run long, a model works in episodes. Each episode accumulates context until the window overflows and the transcript is compacted or dropped; the next episode then reconstructs where it was — re-reading the same files, re-deriving decisions it already made, re-discovering dead ends it already ruled out. This context-recovery cost is pure overhead: activity with near-zero decision efficiency. It is not a prompt-quality problem that a better instruction removes; it is a structural property of unbounded session growth. The longer the horizon, the larger the fraction of elapsed time the system spends paying this tax instead of advancing the frontier.

3.3 Unbounded context is not institutional memory

The obvious response — never drop context, just make the window bigger — solves the wrong problem. An ever-growing transcript is not memory. Institutional memory is curated, addressable, and reusable: a specific technique that worked, a baseline that was rejected and why, a verifier that caught a fake result. A monolithic transcript is none of these; it is an undifferentiated log in which the few reusable facts are buried among the reasoning that produced them, and it is precisely what lossy compaction shreds first when the window finally overflows. Durable capability has to live *outside* the model, in persisted state that can be written deliberately, retrieved selectively, and carried into a future run.

3.4 Process data disappears when only the final artifact survives

Most of the value of a research run is in its process, not its final artifact. The manuscript or merged patch records the destination; it discards the map of how the work got there — what was attempted and failed, which measurement was found contaminated, why one design was chosen over another. When a system keeps only the final artifact, that process knowledge is thrown away at the end of every run, so the next run cannot stand on it and is free to repeat the same mistakes. A system that is meant to compound must persist the trajectory and the verdicts, not just the deliverable, and must treat those process records as first-class, auditable state.

3.5 Design requirements for compounding research

These failure modes are not independent bugs; together they define what a long-horizon research system must provide. Five requirements follow directly, and each is met by a concrete mechanism in the architecture that Act II documents (Section 5).

- R1. Persistent objective.** A durable, operator-anchored objective that survives restarts and upgrades, so the system always knows what it is pursuing and never silently re-plans a campaign

from a stale premise.

- R2. Role separation.** Distinct model-driven roles for decision, execution, and verification, so no single stateless call owns all three and completion is never a self-report from the role that did the work.
- R3. Bounded context.** Per-session context held to an explicit bound and re-seeded from curated working memory, so the runaway compaction that drives context-recovery cost cannot occur.
- R4. Independent evidence.** Progress and completion judged by an authority that inspects artifacts and execution logs against evidence, rather than inferred from activity, filenames, or token volume.
- R5. Reusable runtime state.** Verified evidence deposited as durable, addressable capability — memory, skills, tools, verifiers, routing, evaluations — that the next unknown task inherits instead of starting from zero.

Act II turns these five requirements into a running system: a three-plane architecture that separates decision authority, execution, and evidence, and the runtime that keeps them coupled over a persisted project.

ACT II · THE MACHINE THAT TURNS WORK INTO CAPABILITY

Continuity requires an institution, not a longer prompt.

Argus organizes persistent objectives, bounded role contexts, real execution, independent review, and durable evidence into one research runtime.

4 Argus Life and the Technical Spine

Design objective

Turn the report's master spine — an unknown objective entering a dense-intelligence runtime, passing an evidence gate, updating durable runtime state, and meeting the next task from a higher floor — into named, source-grounded components. Argus Life keeps the lifetime objective moving; each mission is a bounded unit with a first-class outcome.

Act I (Section 3) derived five requirements for a compounding research system: a persistent objective, role separation, bounded context, independent evidence, and reusable runtime state. Act II is the running system that meets them. This section is the map from the expanding-frontier spine (Figure 1) to the implementation — which component owns each stage and what durable artifact it leaves behind — and the sections that follow expand each stage: the four roles and three planes (Section 5), the mission lifecycle and durable state (Section 6), the evidence and review architecture (Section 7), and the reliability and resource machinery (Section 8).

4.1 Life keeps the objective moving; missions stay bounded

Argus Life is the persistent context of a campaign: a durable, operator-anchored objective and its committed configuration, held on disk and recovered across restarts and upgrades. A **STANDING** objective is persisted atomically to `continuous.json` and is never silently re-decided (Section 6). Life itself issues no provider work; it exists so a long campaign has a stable identity. The unit of *work* is a mission: one bounded Engineer↔Reviewer run that always ends in a first-class terminal or paused status, never an unbounded spin. Life supplies continuity; missions supply bounded, auditable progress. That division is what lets the system run for days while every increment stays a discrete, inspectable event.

4.2 The spine, mapped to owners and durable outputs

Table 3 maps each stage of the master spine to the component that implements it and the durable artifact it writes. The mapping is deliberately one-to-one: every conceptual stage has exactly one implemented owner and one place its result is persisted, so a reader can move from the figure to the code without guessing.

Two properties of the mapping matter. First, the *evidence gate* is a single role — the Reviewer — not a hardcoded check (Section 5); the authoritative record of a verified achievement is exactly one event on the tape, `research.achievement.certified` (Section 7). Second, *runtime evolution* is not one mechanism but several, and the report is careful not to flatten them into a single arrow.

4.3 Runtime evolution has four distinct owners

The spine's runtime-evolution step — the update $H(t+1) = U(H(t), \text{trajectory}, \text{evidence})$ with the model's parameters held fixed — is realized by four ownership paths, each carrying a different authority.

Master-spine concept	Implemented owner	Durable output
Unknown / OOD objective enters	Manager front door + division (<code>manager/_core.py</code>)	committed vertical + lifetime in <code>continuous.json</code>
Dense-intelligence runtime	LifeSupervisor outer loop (<code>life/supervisor/_core.py</code>) + SkillLoop per mission (<code>loop.py</code>)	<code>events.jsonl</code> tape + project journal
Four persistent roles	Manager, Planner, Engineer, Reviewer	typed role verdicts on the tape
Evidence gate	Reviewer completion authority (<code>reviewer/_core.py</code>)	<code>research.achievement.certified</code>
Runtime evolution	Reviewer <code>skill_ops/wiki_ops</code> ; Scientist distill-on-miss (<code>skills/, wiki/</code>)	versioned skills + immutable wiki sources
Expanded frontier / next task	LifeSupervisor next claim; Planner refill (<code>planner/planner.py</code>)	refilled <code>backlog.jsonl</code> inherited by the next mission

Table 3. The master spine (Figure 1) mapped to implemented owners and their durable outputs. Each stage has one owner and one persistence point.

Automatic.

Some updates are mechanical. On a matcher miss a Scientist distills an immediately usable project-layer skill; after every real Reviewer verdict a *source* wiki RoundCard is written verbatim from the verdict; and a recency memory preamble is rendered before each mission. None requires further judgment (Section 6).

Model-proposed.

The Engineer *proposes* — it drafts new skill content and, at a turn boundary, emits a curated working-memory handoff. A proposal is a candidate, not yet a committed change.

Reviewer-authored.

Durable knowledge is *committed* by the Reviewer, the sole author of `skill_ops` and `wiki_ops` and the validator of the canonical `checkpoint.json`. A synthesized wiki page is admitted only when its evidence spans verify verbatim against a source (Section 7).

Operator-invoked.

Some evolution is human. An operator edits the machine identity, exports or overwrites skills, sets configuration knobs, switches a role’s backend through the Manager, or leases GPUs — explicit actions, never inferred by the runtime (Section 8).

Keeping these four distinct is a design commitment: the system never presents an automatic distillation as if a human authorized it, nor a Reviewer-certified skill as if it were an unchecked model suggestion. Authority over the runtime is always attributable.

5 Four Roles, Three Planes, One Research Runtime

Design objective

Separate authority two ways at once. Three planes physically divide decision, execution, and evidence, so judgment and record-keeping can never be confused. Four persistent roles carry every scientific and scheduling decision across a narrow typed interface — a Python dataclass with named fields — so no role can quietly reach past its contract.

Argus is one research runtime organized as three cooperating planes and driven by four persistent roles. The **control plane** holds the roles and scheduler that decide *what to run and whether it is done*; the **execution plane** runs *one mission at a time* against real tools, providers, and budgets;

and the **evidence plane** records the run on an append-only tape and never decides it. Figure 3 shows the planes and their interfaces, and Figure 4 is the accepted schematic with the four roles inside the persistent runtime. Table 4 assigns each plane its authority.




Figure 3. The three-plane architecture. Control-plane roles dispatch work into the execution plane, which emits typed events into the evidence plane; the evidence plane is read back by both other planes but holds no decision authority. Drawn deterministically from the module map (Appendix D).

The end-to-end control path is a single spine: CLI → `Manager.divide` → runtime wiring → `LifeSupervisor` → `SkillLoop` → `SupervisedEngineer` ↔ `Reviewer-until-done`, resting on the core services, the single anti-stall mechanism, the pluggable verticals, and the provider backend. Every provider call, regardless of role, is issued through one application-level gateway (`core/run_gateway.py`) that assigns a `call_id` and stamps timing — the single interception point that makes tracing, budgeting, and the per-call audit of Section 7 possible. The concrete backends (Codex, Claude Code, Copilot CLI) sit behind an adapter; a deterministic in-memory backend drives tests. The four roles below each emit a typed decision interface, so authority is explicit and verdicts are machine-checkable. Three form a numbered execution spine (Planner L4, Engineer L1, Reviewer L2); the Manager spans the whole pipeline and sits outside the numbering. (A fifth role, the Curator, runs only in team/subagent mode and is out of scope here.)

5.1 Manager — front door and stage authority

The Manager converts operator free text into a committed campaign in a single grounded model call that emits three structured axes at once — `CONFIG` (e.g. a backend/model change), `CONTROL` (whether to dispatch at all), and `ROUTE` (chat vs. task and which vertical). A `CONTROL: NO_DISPATCH` decision is authoritative: the bridge forces a read-only self-reply and, if that fails, fails closed rather than enqueueing work. A task division yields a `Division`; stage progression is a `StageTransition` (Table 5). The Manager is the sole writer of `current_stage`; the transition’s `action` and `source` fields make both the decision and its justification auditable, including the fail-safe holds that keep

Plane	Primary components	Authority
control	Manager, Planner, LifeSupervisor, daemon (manager/, planner/, life/, daemon/)	decides what to run and whether a project is done
execution	SkillLoop, Engineer, Reviewer, provider gateway, verticals (loop.py, engineer/, reviewer/, core/run_gateway.py)	runs one mission; the Reviewer holds mission-level scientific authority
evidence	event tape, usage ledger, secret guard, provenance, wiki (core/, skills/provenance.py, wiki/)	records and checks; holds no decision authority

Table 4. The three planes and their authority. Control and execution decide; evidence records. Full module map: docs/ARCHITECTURE.md.




Figure 4. System schematic (not a measured performance plot): an operator or objective enters the persistent runtime; the Manager, Planner, Engineer, and Reviewer sit inside it as the four roles that carry every research judgment; and inspectable artifacts and evidence exit back to the operator.

the pipeline coherent when a downstream role is unavailable. The other roles may advise a stage change but never write it.

5.2 Planner (L4) — forward scheduling

When the backlog empties in continuous mode, the Planner proposes the next batch of work as a set of `TaskSpecs` and returns a `PlannerVerdict` (Table 6). Tasks form an optional DAG — each `TaskSpec` has a local `key` and a `deps` list the scheduler maps to real backlog ids — and each carries a structured `scope` threaded to the Reviewer as a backlog tag, so completion authority never re-parses prose to learn a task’s scope. The verdict also has a first-class `waiting` outcome (with a durable `WaitingContract`) for a project correctly blocked on a live external job with no genuinely new high-impact work to queue. The Planner is the sole plan author; the Reviewer may only report that new evidence has overtaken the current plan.

5.3 Engineer (L1) — one round of real work

The Engineer executes a single round against real tools and hardware. Its interface to a provider is the per-call `RunnerOptions` contract, and it receives back a `RunnerResult` (Table 7). Two properties matter for integrity. First, provider-side spend fences (`max_budget_usd`, `max_ai_credits`)

Field	Type / values	Meaning
<i>Division</i>	manager/_core.py	
task	str	the committed objective text
vertical	str	the one selected task shape (never re-classified)
stages	list	the vertical's ordered stage template
<i>StageTransition</i>	manager/_core.py	
action	advance hold rollback complete	the move applied to <code>current_stage</code>
source	manager_llm *_hold	a model decision, or a fail-safe hold

Table 5. Manager decision interface. Only the Manager mutates `current_stage`; a hold writes nothing.

Field	Type	Meaning
<i>TaskSpec</i>	planner/planner.py	
title,	str	human title + full actionable description
objective		
impact_score	int 0–5	parser accepts only high-value work
scope	bounded fi- nal_submission	threaded to the Reviewer as a tag
key, deps	str, list	DAG node id and its dependencies
<i>PlannerVerdict</i>	planner/planner.py	
project_done	bool	project-level completion proposal
new_tasks	list	the next batch
waiting	bool	intentional idle on a live external blocker
restart_daemon	bool	request a controlled daemon restart
checklist_ops	list	per-stage checklist edits (Planner owns it)

Table 6. Planner decision interface (abridged). Scope is structured, never inferred from the objective text.

are populated from an atomic budget *reservation*, not from role configuration, so a role cannot raise its own ceiling (Section 8). Second, the result carries usage-*presence* booleans, so a real zero token count is distinguishable from a provider that simply omitted usage — the evidence plane never invents a number a provider did not report. The Engineer's round-loop reliability guards are documented in Section 8.

5.4 Reviewer (L2) — the completion authority

The Reviewer is the sole authority on completion, and there is no separate hardcoded completion gate anywhere in the system. Each round it inspects the Engineer's output against a structured stage checklist and returns a `ReviewDecision` (Table 8) whose `status` is `done`, `continue`, or `blocked`. Because the Reviewer runs in the same work-tree as the Engineer and is given the mission's execution log with grep recipes, it can audit *how* a result was produced — catching a hardcoded answer, a skipped step, or a faked metric — rather than trusting a summary (Section 7). Its `progress_class` is a structured judgment of forward motion that the infrastructure *counts* but never infers; `skill_ops` and `wiki_ops` are authored here with no Manager gate; and a `final_submission_certified` property is fail-closed, true only when the structured gate is genuinely satisfied.

Why “done” has exactly one owner

Premature acceptance (Section 3) is answered structurally: completion is a single typed verdict from a role that sees the artifacts and the execution log, not a self-report from the role that did the work, and not a keyword the infrastructure fished out of a summary. A prior reviewer path

Field	Meaning
<i>RunnerOptions</i> (per-call knobs)	<code>core/models.py</code>
<code>model, reasoning_effort</code>	provider model and effort level
<code>output_schema_path</code>	structured-output schema for the call
<code>sandbox_mode</code>	containment policy for the spawned role
<code>max_budget_usd, max_ai_credits</code>	spend fences from a reservation, not config
<i>RunnerResult</i> (per-call outcome)	<code>core/models.py</code>
<code>call_id, call_id_log_correlated</code>	audit correlation to the event tape
<code>*_tokens_present</code>	presence flags: a real zero vs. an omitted usage
<code>cost_usd, pricing_status</code>	usage cost and pricing state

Table 7. Engineer per-call backend contract (abridged). Spend fences come from a reservation; presence flags keep usage accounting honest.

Field	Meaning
<code>status</code>	<code>done continue blocked</code> (the completion verdict)
<code>progress_class</code>	<code>decision evidence setup_only artifact_sync_only none</code>
<code>next_action</code>	concrete next step handed to the subsequent Engineer round
<code>operator_question</code>	the one human-facing block question, if any
<code>scope, checklist</code>	final-submission gate (fail-closed)
<code>skill_ops, wiki_ops</code>	knowledge-system edits (Reviewer is sole author)
<code>backend_unavailable</code>	infrastructure-failure flag, not a judgment

Table 8. Reviewer verdict — the completion contract (`core/models.py`, `reviewer/reviewer_schema.json`).

that scanned the Engineer’s summary with regexes and forcibly rewrote a `done` into `continue` was removed; the constraint now lives in the Reviewer’s own prompt. If the Reviewer is wrong, that is a model-quality limitation (Section 11) — but the authority is never ambiguous, and it is never silently overridden.

6 Mission Lifecycle and Durable State

Design objective

Run one mission at a time, back to back, for days — surviving restarts, upgrades, and provider outages — without ever re-deciding a committed campaign or killing a mission mid-flight. Keep the durable truth on disk: the append-only event tape is the system of record, and every other file is a derivable view. Provider sessions are bounded caches, not the sole memory.

The runtime is the 7×24 loop that turns a durable objective into a stream of missions. Figure 5 shows the lifecycle and its recovery edges. This section describes the states, the process model, the restart/resume semantics, and the on-disk state that together make a long campaign durable.

6.1 From backlog claim to mission outcome

The `LifeSupervisor` runs a synchronous outer loop: it claims the next backlog item, injects a non-authoritative recency memory preamble without polluting the raw objective, runs the mission




Figure 5. Mission lifecycle and recovery edges. The primary spine (indigo) claims a backlog item, runs an Engineer↔Reviewer mission, and acts on the Reviewer verdict; the plan-next loop (grey) refills the backlog; recovery edges (terracotta) route paused and blocked missions to first-class statuses and back. The daemon boundary governs draining, handoff, and identity-checked resume.

through the execution plane, accounts the cost against budget gates, writes a journal entry, and — when the backlog empties in continuous mode — invokes the Planner for the next batch. The backlog claim (*pending*→*running*) is atomic, so two workers can never take the same item; orphaned *running* items from a prior crash are reaped on startup. The loop is deliberately synchronous — one mission at a time — which is what makes the mission boundary a clean, well-defined point for stopping, upgrading, and resuming.

Every mission ends in a first-class terminal or paused status, never an unbounded spin. Terminal outcomes include *done*, *blocked*, *max_rounds*, *no_progress*, and *error*; paused outcomes include *paused_budget*, *paused_provider_cooldown*, *paused_provider_fence*, and *infra_blocked* (full taxonomy in Appendix A). A paused mission is not a failure — it records that the campaign is correctly waiting on an external condition. Separately, a *replan_requested* status is a *control transfer* back to the Planner, not a successful completion and not a *done*/*failed*/*blocked* outcome: the current item returns to *pending* and then passes the normal Planner rate and budget gate before any re-plan takes effect.

6.2 Persisted identity and safe resume

For unattended operation the runtime runs as a detached daemon worker (`daemon/life_worker.py`, roughly 1,800 lines) wrapping the supervisor. On `SIGTERM` or `SIGINT` it drains to the next mission-tick boundary rather than interrupting a live mission, so an in-progress evaluation or a result being recorded is never severed. A run is protected by both the atomic backlog claim and a per-project `daemon.pid` lock, so a second daemon cannot drive the same project.

A 7×24 system must be upgradable without losing a campaign. Argus handles a source change with a **blue/green self-handoff**: the worker detects a change in its own source signature, spawns a candidate child under a handoff protocol with a bounded generation count and a minimum inter-

Path (per project)	Owner	Content	Mutability
events.jsonl	evidence plane	canonical typed-event timeline	append-only
backlog.jsonl	scheduler	pending / running / done work items	atomic claim
continuous.json	scheduler	STANDING lifetime + campaign config	rewrite
inbox.jsonl	operator	queued operator messages / answers	append + drain
checkpoint.json	Reviewer	hard-capped curated working memory	atomic rewrite
daemon.pid	daemon	concurrency lock	exclusive
daemon.status.json	daemon	liveness / status snapshot	rewrite
identity.md [†]	global	author / machine identity	operator
skills/ [†]	global	shared skill library	versioned
config.json [†]	global	operator knob overrides	operator

Table 9. Principal persisted state (abridged; full map in Appendix B). The event tape is the only append-only source of truth; everything else is derivable or resettable. [†] under the global root `~/.argus-skill/`.

handoff interval, and hands the campaign over at a clean boundary (`daemon/handoff.py`). The Manager-to-daemon handoff *identity* (objective hash, vertical, lineage generation) is persisted and recoverable from the immutable event tape. Resume is identity-checked: a `-resume-continuous` start adopts only a persisted campaign whose identity matches; a mismatched or missing identity forces a real Manager division rather than blindly re-attaching. Crash and reboot restart is owned by the `systemd -user` unit (Section 8) plus `loginctl enable-linger`, not bespoke keep-alive logic.

6.3 Bounded sessions and curated working memory

Argus does *not* depend on one unbounded provider transcript. The Engineer and Reviewer use separate resumable sessions, so recent context and provider prefix caches stay useful, while a structured `checkpoint.json` in the project state directory carries curated cross-round state (`engineer/checkpoint.py`, `apps/_runtime.py`). Its `CheckpointState` fields cover the goal, verified work, dead ends, one open blocker, the next step, and bounded environment facts. The *Reviewer* authors it: the Engineer ends a turn with a concise handoff proposal, and the Reviewer validates that proposal against the artifacts and writes the canonical checkpoint by atomic rewrite. Python enforces hard item and character caps — not only in the prompt or schema — so the checkpoint stays a *current-state* object rather than an append-only log; stale detail is deleted while ground truth remains in on-disk artifacts.

The resume window is explicitly bounded. By default the runtime rolls each role’s thread after three rounds on that thread, or when the preceding call reports at least 1.5 million input tokens; the next call opens a new session and the Engineer is reseeded from the checkpoint (`engineer/runner.py`, `core/knobs.py`). If provider-side compaction occurs before a roll, the next round resends the full static preamble. This combines short-window context reuse with a structural limit on transcript growth, and it fails soft: a missing or garbled checkpoint keeps the last good one rather than clearing memory.

6.4 The on-disk state map

Core path helpers define the global and per-project roots (`core/paths.py`); runtime composition derives the checkpoint path inside the project state directory (`apps/_runtime.py`). A global root `~/.argus-skill/` (overridable via `ARGUS_SKILL_HOME`) holds cross-project identity, shared skills, and an operator knob store; each project lives under `projects/<fingerprint>/`. Table 9 lists the principal state files, their owner, and their mutability. The canonical timeline is the append-only `events.jsonl`; every other file is a derived view or working state that can be rebuilt or safely reset. Long-term knowledge lives in two systems on this state. Skills are versioned Markdown files distilled

on a matcher miss and revised from real usage (`skills/`); a per-project wiki records an immutable *source* RoundCard after each Reviewer verdict, from which the Reviewer may later synthesize evidence-cited pages (Section 7). Both compound capability across missions without becoming an unbounded prompt.

7 Evidence and Independent Review

Design objective

Make every reported number reconstructable and every completion auditable. The evidence plane records the run on an append-only tape, lets the Reviewer audit how each result was produced, redacts credentials before persistence, hashes published artifacts for provenance, and labels every quantity by the strength of its evidence. A verdict may be wrong, but it can never be unfalsifiable.

7.1 The immutable event tape

The canonical record of a project is an append-only event tape, `events.jsonl`, written through typed events from a catalog of 112 event types across 11 categories, of which 75 carry a validated payload schema (`core/event_catalog.py`; Appendix C). Within it exactly one event, `research.achievement.certified`, is the authoritative record of a verified achievement, emitted only on a `done` verdict carrying the structured achievement payload (Section 5). A campaign's decisions, spend, and outcomes can be replayed from the tape after the fact, which is what makes an audit possible at all.

7.2 Execution-log audit and call correlation

Trusting a self-report would defeat the completion contract, so the Reviewer checks the Engineer's work directly. The two roles share a work-tree, and the Reviewer's prompt includes an audit section pointing at the mission's own `events.jsonl` with grep recipes (`engineer/runner.py`), so the Reviewer can inspect *how* a result was produced — detecting a hardcoded answer, a skipped step, a cheat method, or a faked metric — rather than accepting the summary. Each provider call is correlated to the tape by `call_id`, and usage-presence booleans keep a real zero distinct from a missing measurement (Section 5). This verifies the number is honest, not that the science is good.

7.3 Anti-cheat by construction

For benchmarks whose metric could be gamed by memorizing a fixed evaluation input distribution, the evaluation input is randomized per run, so a solution cannot hardcode the statistics of a known distribution — a fast-because-faked kernel is designed to be impossible to pass off as genuinely fast. Combined with real public benchmark data and the per-round audit trail, this closes the metric-exploitation failure mode by construction rather than by after-the-fact inspection. Integrity is per-benchmark and remains an ongoing surface: a new benchmark can present a new exploit the current checks were not written for (Section 11).

7.4 Credential redaction before persistence

A task-agnostic secret guard (`core/secret_guard.py`) scrubs credentials from newly written artifacts before they reach the event log or the Reviewer's prompt, and redacts persisted event and usage copies, matching authorization and API-key patterns, provider tokens, and generic key/value secrets. Redaction never judges scientific quality; when it rewrites an artifact it emits `round.secret_redacted` so the Reviewer can rebuild the affected provenance hashes, and a scan error or an oversized artifact explicitly *blocks* unconditional certification rather than silently passing — the same fail-closed stance as the completion gate.

7.5 Provenance, snapshots, and digests

Published evidence carries provenance. For website-sourced results, the retrieval records an HTTP status and a raw-HTML SHA-256 for each source page; for corroborated results, the bundle identifies an artifact in a named external project by logical project name, artifact id, and SHA-256 — with no local absolute paths, usernames, or session ids. Table 10 separates a `website_snapshot` — a result quoted verbatim from the live public record — from a result that additionally carries committed external-project artifact-digest metadata. The artifact bytes are not committed in this repository (`technical_report/evidence/README.md`). Result *signing* is deliberately not claimed: no cryptographic result-signing component exists in the current tree, and the report does not imply one.

Corroboration category	What it means	What it may claim
<code>local_artifact</code>	A public result with a committed artifact-digest record for a named external project (logical project + artifact id + SHA-256).	An Argus result for that run, on stated hardware.
<code>website_snapshot</code>	A result card quoted verbatim from the live public record, with no artifact-digest corroboration yet.	A scoped public claim under its stated protocol.
External baseline	A published third-party or human target used for comparison.	A reference target; explicitly <i>not</i> an Argus result.
Ongoing / no claim	A program running without certified, reproducible completion.	Status “ongoing”; no numeric Argus claim.

Table 10. Evidence categories. All six public results originate as `website_snapshot`; the bundle’s separate corroboration field marks two as `local_artifact` because artifact-digest metadata is available (Section 9).

7.6 Completion is a single certified verdict

The measurement-integrity architecture converges on one rule: a program is complete only when the Reviewer certifies it against the full-pipeline stage checklist, with the `final_submission_certified` property fail-closed. There is no hardcoded completion gate and no independent validator behind the Reviewer; an earlier family of roughly twenty machine “validate-*” quality gates was removed once completion authority moved to the Reviewer’s checklist, so completion has exactly one owner and one auditable record. This concentrates authority in a single decision point — examined honestly in Section 11, since the Reviewer is a model with no independent oracle behind it — but it also means completion is never the product of two components silently disagreeing.

8 Reliability, Resources, and Operation

Design objective

Treat execution drift — not weak reasoning — as the primary risk of a multi-hour run. Every degradation mode maps to a named guard that fails soft and is precise about whether it may interrupt a live call. Route every provider call through one priced, reservation-backed gateway so spend is bounded and attributable. Give the operator a live, authenticated cockpit over the same tape, and deploy as a standard service that drains cleanly.

The reliability model is a set of guards that keep a mission bounded and honest. Figure 6 sketches how they wrap one Engineer/Reviewer loop over a persistent budget, event-log, and artifact rail. None of these guards decides whether the research is good: each enforces a task-agnostic invariant and fails soft — reliability is added *around* the agent’s judgment, never in place of it.




Figure 6. System schematic (not a measured performance plot) of the long-running reliability cycle: Planner, Engineer, and Reviewer run one loop; the Reviewer maintains a bounded curated checkpoint and a decision-progress class that guides the next Engineer round; supervised background jobs run alongside; and a safe round boundary returns a stalled or budget-exhausted mission to the Planner.

8.1 Decision progress, not activity

To separate real progress from busy churn, the Reviewer labels each round with a `progress_class`: `decision` or `evidence` for genuine forward motion, or `setup_only`, `artifact_sync_only`, or `none` for nondecision work. The runtime *counts* these structured labels; it never infers progress from tool activity, filenames, or token emission. Two consecutive nondecision rounds trip the stall guard, and a 1,800-second decision-progress budget — measured from the last decision or evidence increment and checked only between rounds — ends a mission that is spinning without ever cutting useful in-flight work. A separate soft/hard round escalation (12/24) catches the rarer case of a mission that makes evidence progress every round but never passes its gate, asking the Reviewer to return `blocked` and, failing that, force-ending so the Planner re-plans. Table 11 is the full guard matrix.

Guard	Trigger / action	Default	Interrupts live call?
Backend-failure recovery	failure streak → backoff, then end as error	2 / 15 s	no (round boundary)
No-progress bail	consecutive rounds with no engineer message	2	no
Decision stall	consecutive nondecision <code>progress_class</code> rounds → end	2	no
Decision-progress budget	seconds since last decision/evidence → end	1,800 s	no
Soft round escalation	round index → ask Reviewer to return <code>blocked</code>	12	no
Hard round escalation	round index → force-end as <code>blocked</code>	24	no
Effective-progress watchdog	no file change / non-token event → kill subprocess	3,600 s	yes
Runner hard-idle	no runner activity → interrupt round	3,600 s	yes
In-round compaction limit	auto-compactions within one round → end round	3	yes

Table 11. Reliability guards (`engineer/runner.py`). Round-boundary guards never sever a live call; the three in-round guards fire only against a subprocess emitting token noise or looping without effective progress. All are env-overridable and disable at 0.

8.2 In-round watchdogs and backend failure

Three guards operate *inside* a round because they target a subprocess that has stopped doing useful work. The effective-progress watchdog kills a provider subprocess that keeps emitting heartbeat or token noise for 3,600 seconds with no file change and no non-token session event; the runner hard-idle watchdog interrupts a round with no runner activity at all; and the in-round compaction guard ends a round whose session auto-compacts three times. Each default is intentionally long — research turns legitimately spend many minutes reading, planning, or waiting on model-side recovery — and each disables at 0. A backend failure streak triggers a bounded backoff and, if it persists, ends the mission as an error. A degraded reviewer backend is routed to the structured `backend_unavailable` status — treated as infrastructure and never as a `continue` — after a historical blind-gate incident in which a stale import-time schema path ran the completion gate blind for roughly ninety minutes; the reviewer now preflights its schema file and raises if it is unreadable rather than proceeding blind.

8.3 One gateway, pricing, and budget reservations

All backend execution flows through a single application-level gateway, `RunExecGateway.execute` (`core/run_gateway.py`). Token usage is priced by a small module (`core/pricing.py`, default model `gpt-5.5`) and accumulated in a *presence-aware* usage ledger (`core/usage.py`) that distinguishes a genuine zero from omitted usage, so no cost is ever fabricated from missing data. Spend control is reservation-based and event-backed: before a call runs the runtime creates a budget reservation, recorded as one of `budget.reservation.created`, `.denied`, `.settled`, or `.released`, with `budget.unpriced.blocked` and `budget.fence_breach.blocked` for the unpriceable or fence-breaching cases. This reservation-then-reconciliation cycle — settle on completion, release on abort — keeps the ledger honest, and the provider-side spend fences on a call (`max_budget_usd`, `max_ai_credits`) are injected from the reservation, not from role configuration, so a role cannot raise its own ceiling. The principal live defaults are a per-mission preflight cap of \$30, a daily cap of \$180, an optional global daily cap (off unless set), and at most two active daemons across projects. Beyond dollars, a provider quota permit (`core/provider_quota.py`) turns a quota denial into a paused status rather than an error.

8.4 Background jobs and the cadence wait

Not all work is short. When the Engineer starts a long experiment — for example a GPU training run — it launches a supervised background job that has its own monitor: the monitor checks health on an interval, early-stops on failure, and reports a terminal result to the operator inbox. The Engineer reads a registry of in-flight jobs (`engineer/background_subagents.py`) and classifies each as self-watched or needs-attention; when it explicitly replies `WAIT_FOR_SUBAGENT` for a healthy self-watched job, the runner yields to that job’s own monitor cadence instead of busy-polling the Reviewer. It is an agent-led choice to wait, not a harness decision, and a stale or misfired sentinel simply falls back to a normal reviewed round rather than hanging the loop.

8.5 Coordinating scarce hardware: the GPU lease

On managed GPU hosts a scheduler reclaims cards that sit idle, so an operator runs a low-duty “keep-alive” loader to hold them. Argus ships an explicit *lease* tool for this (`argus_skill/tools/gpu_lease.py`): the recommended path, `gpu_lease run -- <cmd>`, frees the cards, records a lease for the job’s lifetime, runs the command, and on exit re-parks the keep-alive only if no other lease is active. A `-detach` mode lets a supervisor own the lease independently of the mission, and process termination only ever targets the configured match token.

Scope of the GPU lease

The GPU lease is a real, robust *operator/agent-invoked* capability, not an automatic runtime service: nothing in the mission scheduler or engineer loop calls it implicitly. It is resource-

governance tooling the agent may choose to use — which is exactly what the live integration supports, no more.

8.6 Operator cockpit and observability

Argus is not only a daemon; it is a live workbench over the same event tape the system runs on. The cockpit has a terminal surface (an Ink TUI), a web surface, and an HTTP/WebSocket API (`webapi/server.py`); there is no embedded Python REPL. The front door routes free text through the Manager, so an operator can say “switch to the Copilot backend” in plain language; the four roles are shown as peers and a mission is followed round by round via `-status`, `-watch`, and `-follow`, backed by the typed event stream. Because the Engineer and Reviewer share a work-tree and every round is persisted, the operator can open exactly what was produced and read the Reviewer’s reasoning rather than a summary. The decision interface to a human is deliberately narrow: at most one blocking `operator_question` per mission, queued to and drained from an operator inbox. Most inspection endpoints are open; commands, the live WebSocket, and global metrics and trash views pass through an authentication dependency and require an `Authorization` token when one is configured — routine inspection stays friction-free while every mutation and sensitive view is gated.

8.7 Deployment, containment, and cost

The intended deployment is a per-project daemon supervised by `systemd -user` (`deploy/argus-skill.service`): it starts the daemon in the foreground so the service manager owns the lifecycle, drains to the mission boundary on stop (`TimeoutStopSec 1800`, `SIGTERM`, no mid-mission `SIGKILL`), restarts on failure with a bounded restart budget, and is `cwd`-bound to a project fingerprint; `loginctl enable-linger` carries it across logout and reboot. A daemon replacement or upgrade drains to the same clean boundary rather than killing a live mission, and the identity-checked resume (Section 6) ensures the restart continues the right campaign. Each builder role can run under a provider sandbox (`core/sandbox.py`) confined to the project working directory; the policy *refuses* a fake sandbox that would present as confined while remaining writable, and a sandboxed builder may never write the global state root or the provider’s own configuration. Spend is bounded by the reservation system above — defaults, not guarantees: a poorly bounded campaign can still be expensive, and 7×24 operation is not free (Section 11).

9 Public Results Across Six Arenas

Design objective

Report each public result as a scoped claim under the exact protocol it was measured under, using human or paper-reported references first where available, with its evidence status attached. Units differ across arenas and are never cross-normalized; a website snapshot is never presented as a reproduced artifact.

Argus publishes a live public record [1]. This section reproduces its six current arena results. Figure 7 presents them as small multiples — one panel per arena, each with its own unit — because the arenas measure different quantities (kernel rank, bits-per-byte, wall-clock seconds, solve rate, a gap score) and any shared axis would be misleading. Table 12 gives the exact protocol, result, published comparison, and evidence status for each row.

9.1 Reading the arenas

Each number means exactly what its protocol row says and nothing more. On **SOL-ExecBench** [6] the headline is a *global rank* (#6) against all entrants on 101 B200 kernels, with two #1 finishes and seven top-3 placements; the two head-to-head wins over Recursive [7] are the site’s note, not the framing. **nanochat** [5] is reported on both B200 (0.9636 BPB against a 0.9646 human SOTA)

Public results across six arenas (units differ; panels are not cross-normalized)



Figure 7. Public results across six arenas, drawn deterministically from `technical_report/evidence/website_results.json`. Panels are not cross-normalized; each carries its own unit and evidence status. Direction is shown where defined; Arbor retains the site’s reported metric without inferring its direction. Argus bars are indigo; human/system baselines are lighter.

and H100 (0.9855 against 0.9879), each under a 5-minute training budget — lower bits-per-byte is better. The **nanoGPT speedrun** reproduces a fast GPT-training loop [4] on 8×H100 and reports 79.77s over $N=10$ against the same-device human record of 80.18s. **AARRI-Bench** reports 63 of 82 research-intern tasks (76.8%) against a paper-reported best of 68.3%. For **Arbor**, the website reports a “gap” value of 28.0 alongside Arbor 20.83, Claude Code 8.33, and Codex 6.25. The snapshot does not define the metric’s direction, so this report quotes the published values without interpreting the ordering.

9.2 Artifact-digest corroboration versus website snapshot

Two of the six rows carry digest corroboration, and the difference is stated rather than smoothed over. For the **nanoGPT** 79.77s row, the bundle references an $N=10$ verifier-certified artifact in the external `nanogpt-speedrun-h100` project — the verifier line records `valid=true`, `n=10`, a validation loss of 3.2776, a one-sided $p(\text{mean} < 3.28) = 0.0040$, a train time of 79.77 ± 0.06 s, and a `seal=ok`. For the **nanochat B200** 0.9636 row, the bundle records frozen-scorer, single-seed evidence of `MEAN_VAL_BPB = 0.963634` with a frozen-scorer exit code of 0 and a genuine SM100 flash-attention kernel, in the `nanochat-mission-b200` project. This repository commits the logical artifact IDs, verifier excerpts, and SHA-256 digests, not the external artifact bytes. The other four rows — SOL-ExecBench, nanochat H100, AARRI-Bench, and Arbor — are `website_snapshot` with no artifact-digest corroboration yet, linked to the machine-readable bundle `technical_report/evidence/website_results.json`. No corroboration file was fabricated where artifact-digest metadata was not available.

Arena	Protocol	Argus result	Published comparison	Status
NVIDIA SOL-ExecBench	B200; 101 kernels	Global #6; 2×#1; 7 top-3	Global rank vs. all entrants (headline)	snapshot
nanochat (B200)	5 min; 1×B200; 426 attempts	0.9636 BPB	Human SOTA 0.9646	digest
nanochat (H100)	5 min; 1×H100; 37 mechanisms	0.9855 BPB	Human SOTA 0.9879	snapshot
nanoGPT speedrun	8×H100; $N=10$	79.77 s	Same-device human #83: 80.18 s	digest
AARRI-Bench	82 research-intern tasks	63/82 (76.8%)	Paper-reported best 68.3%	snapshot
Arbor (RUC NLPir)	Math-reasoning data	28.0 gap	Arbor 20.83; Claude Code 8.33; Codex 6.25	snapshot

Table 12. The six public results [1], verbatim from the evidence bundle. Human and paper-reported references are primary where available; the SOL-ExecBench headline is its global rank. SOL-ExecBench and Arbor also preserve the site’s published system/agent comparisons. “digest” = committed artifact metadata; “snapshot” = public website snapshot.

What these results do and do not claim

Each row is a real, scoped claim under its stated protocol and evidence tier. The table does *not* claim universal state of the art, correctness outside these protocols, or that a snapshot is a reproduced artifact. Converting the four snapshot rows into artifact-digest records tied to a named commit is an explicit roadmap item (Section 12).

10 Research Portfolio: 41 Papers Across Six Programs

Design objective

Report the research output as a de-duplicated inventory, broken down by program and by manuscript/draft status, compared only to human-authored literature and baselines. Count what exists; claim no acceptances.

Beyond the benchmark arenas, Argus runs an optional idea-to-submission research pipeline (the **research** vertical), which drives literature grounding, experiment design, execution, results analysis, and manuscript writing through per-stage checklists and a Reviewer peer-review gate. Its public output is a collection of **41 papers** [2]: 35 manuscripts and 6 drafts, with duplicate versions removed, across six research programs. Figure 8 and Table 13 give the breakdown.

Program	Manuscripts	Drafts	Total
Multimodal & Vision-Language Models	13	3	16
Cognitive Bias in LLMs	9	0	9
Efficiency, Compression & Decoding	6	1	7
LLM Agent Methods	3	2	5
World Models	2	0	2
State Trace & Auditability	2	0	2
Total	35	6	41

Table 13. Research portfolio by program and status [2]. The machine-readable inventory is `technical_report/evidence/paper_inventory.json`.

Research portfolio — 41 papers across six programs

35 manuscripts + 6 drafts · human-authored baselines only · de-duplicated inventory, not accepted papers




Figure 8. The 41-paper portfolio by program and status, drawn deterministically from `technical_report/evidence/paper_inventory.json`. Manuscripts are indigo, drafts lighter. The count is a de-duplicated inventory compared only to human-authored literature/baselines — not a count of accepted papers.

10.1 The six programs

The portfolio concentrates in a few areas. **Multimodal & Vision-Language Models** is the largest program (16 papers), followed by **Cognitive Bias in LLMs** (9), which studies train-free mitigation of biases such as anchoring, base-rate neglect, and status-quo defaults. **Efficiency, Compression & Decoding** (7) covers inference-time methods; **LLM Agent Methods** (5) covers agent techniques; and two smaller programs, **World Models** (2) and **State Trace & Auditability** (2), round out the set. Each paper is positioned against human-authored literature and baselines; the human comparison, where the site provides one, is preserved in the inventory.

10.2 What the count means

Three qualifications keep the number honest. First, 41 is a *de-duplicated inventory* of research runs that produced a manuscript or draft; it is not a count of accepted papers, and this report makes no acceptance claim. Second, the manuscript/draft split (35/6) is reported as-is: a draft is an in-progress artifact, not a finished manuscript. Third, the collection is compared *only* to human-authored literature and baselines — paper count and paper quality are never compared to any other agent or model. The research pipeline is a capability that produces auditable manuscripts under the same completion authority as every other mission (Section 5); it is not a claim that any of those manuscripts has cleared external peer review.

11 Limitations and Open Engineering Problems

Argus gives an agent real judgment, real tool access, and real persistence. It does not guarantee good science, and a report that documented reliability mechanisms without documenting their boundaries would itself be misleading. The limitations below are material, and several are open engineering problems the roadmap (Section 12) targets directly.

11.1 Model and objective ceilings

Provider dependence.

Argus drives an external agent CLI (Codex, Claude Code, or Copilot CLI) and inherits that provider's capability, availability, latency, and pricing. An outage, a regression, or a rate limit is felt directly, and the quality ceiling of any run is the underlying model's.

Objective quality is the operator's.

The system optimizes the objective it is given. A vague or trivially gameable objective yields vague or gamed work; Argus enforces integrity of *measurement*, not wisdom of *goal*.

No guarantee of novelty, correctness, or SOTA.

Nothing in the runtime guarantees that an idea is novel, that a proof or result is correct, or that a number is state of the art. Those are properties of the agent's judgment and the problem's difficulty, and the system deliberately does not fake them.

11.2 Sole-reviewer completion authority

The Reviewer is the single completion authority, and it is a model. A wrong Reviewer can accept an incomplete result or reject a finished one. Argus reduces this risk by giving the Reviewer the shared work-tree, the execution-log audit, and a structured checklist (Section 7), and by routing a degraded reviewer backend to an explicit infrastructure-failure status rather than a `continue`. But it cannot eliminate model error, and there is no independent oracle behind the Reviewer. The historical blind-gate incident (Section 8), in which a stale schema path ran the completion gate blind for roughly ninety minutes, is the concrete failure mode of concentrating authority in one component; the fix hardened the transport, not the underlying concentration of authority, which remains an open problem.

11.3 Evidence scope and integrity coverage

Four of six results are snapshots.

Only two of the six public results (Section 9) carry committed digest records referencing artifacts in named external projects; the artifact bytes are not in this repository. The other four are public website snapshots linked to the evidence bundle. They are real, scoped claims under their protocols, but they are not yet independently reproducible from a named commit, and the report does not present them as if they were.

Integrity is per-benchmark and ongoing.

The anti-cheat construction is task-agnostic, but a determined agent explores the exploit surface of each specific benchmark. A new benchmark can present a new exploit — a leaked test set, an unrandomized input, a side channel — that the current checks were not written for. Measurement integrity is an ongoing engineering surface, not a solved problem.

Paper inventory is not acceptances.

The 41-paper portfolio (Section 10) is a de-duplicated inventory compared to human-authored literature; it is not a count of accepted papers, and no acceptance is claimed.

11.4 Operational limitations

Cost.

Long-horizon autonomy consumes provider tokens continuously. Budgets and reservations bound spend, but 7×24 operation is not free.

Secret handling is a scanner, not a proof.

The credential guard scrubs artifacts before they are recorded, but it is a task-agnostic scanner: a scan gap or an oversized artifact blocks certification rather than guaranteeing zero leakage, and operators remain responsible for how secrets enter the environment.

Back-compatibility surface.

Legacy shims (an older global-directory variable and queue helpers) are kept for migration; they are technical debt to retire and add surface an audit must account for.

GPU lease is invoked, not automatic.

The GPU lease (Section 8) is robust but is an operator/agent-invoked tool, not an automatic runtime service; coordinating scarce hardware still depends on the agent choosing to use it.

11.5 Responsible operation

Given these limitations, Argus is meant to be operated and audited, not fired and forgotten. Responsible use means giving it well-posed objectives and honest benchmarks; setting budgets deliberately; reading the Reviewer’s reasoning rather than a headline; treating every number according to its evidence tier (Table 10); and publishing only what the stated evidence supports. The system makes this practical — persisted artifacts, a round-by-round audit trail, a single inspectable completion authority — but the discipline of honest reporting ultimately rests with the operator.

Maturity

Argus is under active development at Technical Report 0.3. The architecture and mechanisms described here are implemented and grounded in the current source tree; the public results (Section 9) are scoped claims under their stated protocols and evidence tiers, not universal-capability guarantees. Treat any performance number according to its tier — an artifact-digest record, a public website snapshot linked to the bundle, or a labeled external baseline — never as a general guarantee of capability.

12 Roadmap

The roadmap follows directly from the limitations in Section 11. Each item is an engineering objective, not a capability promise, and each is stated so that its completion is checkable.

12.1 From snapshot to reproducible evidence

The highest-priority item is to convert the four website-snapshot results of Table 12 — SOL-ExecBench, nanochat H100, AARRI-Bench, and Arbor — into artifact-digest records tied to a named commit, matching the nanoGPT and nanochat-B200 rows that already carry digest corroboration. The target is that every published number can make a specific, auditable claim on stated hardware, verifiable by a third party from a commit and an artifact digest.

12.2 Strengthening measurement integrity

Integrity coverage grows per benchmark: the plan is to extend the anti-cheat construction as each new arena reveals its exploit surface, and to grow the structured stage checklists the Reviewer verifies against — rather than re-introducing a separate family of machine validators that could silently disagree with the completion authority. The constraint is to keep exactly one owner of completion (Section 7) while widening what it is equipped to check.

12.3 Reducing the single-authority risk

The sole-reviewer concentration (Section 11) is the most important open problem. Near-term work reduces Reviewer error through richer execution-log auditing and evidence-grounded verdicts, and hardens detection of a degraded reviewer backend so a blind gate cannot recur. Whether to add independent cross-checks without recreating a hardcoded gate is an open design question rather than a settled plan.

12.4 Retiring debt and lowering cost

The back-compatibility surface (legacy shims and queue helpers) is scheduled for removal once migration paths are confirmed, and the GPU lease is a candidate for tighter but still explicit scheduler coordination. In parallel, tightening budget accounting and the cockpit — better cost attribution, clearer paused/blocked surfacing, lower-overhead observation — makes a 7×24

campaign cheaper to run and easier to steer.

The through-line

Argus today is a research infrastructure for reliable, auditable autonomy at version 0.3. The roadmap is the path from a trustworthy runtime to trustworthy *results* — reproducible evidence, broader integrity coverage, a less concentrated completion authority, and a cheaper campaign — pursued without letting the infrastructure substitute its judgment for the agent’s, or any number outrun its evidence.

A Key Interfaces and Status Taxonomy

The role decision interfaces are summarized in Section 5 (Tables 5–8). This appendix records the status taxonomy they resolve into. Three enumerations govern outcomes: **ReviewStatus** (the verdict), **LoopStatus** (a mission’s terminal or paused outcome), and **PlannerVerdictStatus** (the planner’s event contract), the last carrying an explicit (success, recoverable) policy that determines how the scheduler reacts.

PlannerVerdictStatus	success	recoverable	Meaning
planned	yes	no	a new task batch was produced
completed	yes	no	the project is certified complete
research_incomplete	no	yes	more research is required
paused_budget	no	yes	budget exhausted; recheck later
paused_no_breakthrough	no	yes	no breakthrough this cycle
exhausted_current_methods	no	yes	current methods exhausted
provider_cooldown	no	yes	provider cooldown in effect
provider_fence	no	yes	provider spend fence hit
infra_blocked	no	yes	infrastructure blocker
error	no	no	unrecoverable error

Table 14. **PlannerVerdictStatus** and its (success, recoverable) policy (`core/planner_verdict.py`). Legacy aliases (`done`, `paused_provider_cooldown`, `paused_provider_fence`) map onto these.

ReviewStatus (`core/models.py`): `done`, `continue`, `blocked`, plus the research-pause states `research_incomplete`, `paused_no_breakthrough`, and `exhausted_current_methods`.

LoopStatus (`core/models.py`): the terminal set `done`, `max_rounds`, `blocked`, `no_progress`, `error`, `budget_exhausted`; and the paused set `paused_budget`, `paused_provider_cooldown`, `paused_provider_fence`, `infra_blocked`, `research_incomplete`, `paused_no_breakthrough`, `exhausted_current_methods`; and the control-transfer status `replan_requested`, used only when dynamic planning requests a return to the Planner. Every mission yields one of these statuses — never an unbounded spin (Section 6).

B Persisted State Map

Table 9 (Section 6) lists the principal state files; the root layout is defined in `core/paths.py`, while runtime composition adds the per-project checkpoint path in `apps/_runtime.py`. The global root `~/.argus-skill/` (overridable via `ARGUS_SKILL_HOME`; a legacy `ARGUS_SKILL_LIFE_DIR` shim is retained for migration) holds cross-project identity, the shared skill library, and an operator knob store; each `projects/<fingerprint>/` subtree holds the append-only `events.jsonl`, `backlog.jsonl`, `continuous.json`, `inbox.jsonl`, a `skills/` directory, `daemon.pid`, `daemon.status.json`, and `checkpoint.json`. A runtime path resolver rejects unresolved shell placeholders, and the project-root helper rejects empty, `.`, and `/` fingerprints, so a mis-resolved path fails loudly rather than writing to the wrong place.

C Event Taxonomy

Cross-component state is exchanged through a canonical catalog of **112 typed event types** across **11 categories** (`core/event_catalog.py`); **75** carry a validated payload schema. Table 15 gives the per-category counts and representative events.

Category	Count	Representative events
lifecycle	40	<code>loop.start</code> , <code>round.main.completed</code> , <code>round.review.completed</code> , <code>budget.reservation.*</code> , <code>life.mission.started</code>
skill	20	<code>skill.match.*</code> , <code>scientist.start</code> , <code>skill.outcome</code>
wiki	17	<code>wiki source/page ingest</code> , <code>promotion</code> , <code>validation</code>
planner	9	<code>life.planner.start</code> , <code>life.planner.verdict</code>
research	7	<code>research.achievement.certified</code> (sole authoritative achievement)
agent_io	5	<code>agent.io.start</code> , <code>agent.io.stream</code> , <code>agent.io.complete</code>
daemon	5	<code>daemon start / stop / handoff / status</code>
idea	3	<code>idea lifecycle</code>
provider	3	<code>provider.request.started/completed/denied</code>
usage	2	<code>usage.recorded</code> , <code>cost</code>
operator	1	<code>operator inbox event</code>

Table 15. Event taxonomy: 112 types across 11 categories (`core/event_catalog.py`). Budget, loop, and round events are grouped under lifecycle.

D Evidence Map

Every engineering claim in this report is grounded in the current `argus-skill` repository [3] rather than in transient run directories. Table 16 maps each subsystem to its primary live source.

Subsystem (section)	Primary source path(s)
Three planes (§5)	<code>docs/ARCHITECTURE.md</code> ; <code>core/</code> , <code>manager/</code> , <code>engineer/</code> , <code>reviewer/</code>
Mission lifecycle / daemon (§6)	<code>life/supervisor/_core.py</code> , <code>daemon/life_worker.py</code> , <code>daemon/handoff.py</code>
State & memory (§6)	<code>core/paths.py</code> , <code>apps/_runtime.py</code> , <code>engineer/checkpoint.py</code> , <code>life/memory.py</code> , <code>skills/</code> , <code>wiki/</code>
Role interfaces (§5)	<code>core/models.py</code> , <code>planner/planner.py</code> , <code>manager/_core.py</code> , <code>reviewer/reviewer_schema.json</code>
Execution / budgets (§8)	<code>core/run_gateway.py</code> , <code>core/pricing.py</code> , <code>core/usage.py</code> , <code>core/provider_quota.py</code>
Reliability guards (§8)	<code>engineer/runner.py</code> , <code>regime_jump/</code>
Evidence & integrity (§7)	<code>core/event_catalog.py</code> , <code>core/secret_guard.py</code> , <code>skills/provenance.py</code>
Cockpit / API (§8)	<code>webapi/server.py</code>
Deployment / security (§8)	<code>deploy/argus-skill.service</code> , <code>core/sandbox.py</code> , <code>core/vault-preflight.py</code>
Public results & portfolio (§9–10)	<code>technical_report/evidence/website_results.json</code> , <code>paper-inventory.json</code>

Table 16. Evidence map: subsystem to live source. This report cites repository paths on main, not any transient local filesystem location.

Figure provenance. The four deterministic figures (`system_planes`, `mission_lifecycle`, `public_results`, `paper_portfolio`) are generated deterministically by `technical_report/figures/build-report_figures.py`, with no image-model call. The two result figures read the committed evidence bundles; the architecture and lifecycle schematics encode source-grounded system structure. Reproducible SHA-256 digests are in `figures/REPORT_FIGURES.json`. The two editorial illustration plates

were produced through the image-2 pipeline and carry prompt, inspect, review, and provenance sidecars in `figures/IMAGE2_FIGURES.json`.

E Paper-Inventory Summary

The 41-paper portfolio (Section 10) is enumerated in the bundle `technical_report/evidence/paper_inventory.json`, retrieved from the public research page [2] with an HTTP-200 raw-HTML SHA-256 for provenance. It totals 41 papers — 35 manuscripts and 6 drafts, duplicates removed — across Multimodal & Vision-Language Models (16), Cognitive Bias in LLMs (9), Efficiency, Compression & Decoding (7), LLM Agent Methods (5), World Models (2), and State Trace & Auditability (2), compared only to human-authored literature/baselines; no external acceptance is claimed.

References

- [1] Argus Team. Argus public results. Website snapshot, <https://argusbot.cn/results.html>, 2026.
- [2] Argus Team. Argus research-paper collection. Website snapshot, <https://argusbot.cn/research.html>, 2026.
- [3] Argus Team. argus-skill: A 7×24 system for long-horizon autonomous research. GitHub repository, 2026.
- [4] Andrej Karpathy. nanoGPT: The simplest, fastest repository for training/finetuning medium-sized GPTs. GitHub repository, 2022.
- [5] Andrej Karpathy. nanochat: The best ChatGPT that \$100 can buy. GitHub repository, 2025.
- [6] Anne Ouyang et al. KernelBench: Can LLMs write efficient GPU kernels? *arXiv preprint arXiv:2502.10517*, 2025.
- [7] Recursive. First steps toward automated AI research. GitHub repository, 2026.