


Dense-intelligence research

Continuous proposal—execution—verification over durable state

Invention

Long horizon

Verifiable



Benchmark results and full-suite software evaluation

SWE-Bench Pro · 731 tasks

GPT-5.5/xhigh

59% → 78%

Direct Copilot → Argus 1.41× aggregate tokens

SOL-ExecBench

GPT-5.5 · Codex

Global #6 · 2× #1 · 7 top-3

B200 · 101 kernels

Two head-to-head wins over Recursive

nanochat · B200

GPT-5.5 · Codex

0.9636 BPB

5 min · 1×B200 · 426 attempts

Human SOTA: 0.9646

nanochat · H100

GPT-5.5 · Codex

0.9855 BPB

5 min · 1×H100 · 37 mechanisms

Human SOTA: 0.9879

nanoGPT speedrun

GPT-5.5 · Codex

79.77 seconds

8×H100 · N=10

Same-device human #83: 80.18s

AARRI-Bench

GPT-5.5 · Codex

63/82 · 76.8%

82 research-intern tasks

Paper-reported best: 68.3%

Math-Reasoning Data

GPT-5.5 · Codex

28.0 gap

Math-Reasoning Data

Arbor 20.83 · Claude Code 8.33 · Codex 6.25

41 de-duplicated research artifacts across 6 programs

35 manuscripts · 6 drafts

Architecture is schematic; result cards retain task-native units and are not cross-normalized.

Public arenas: GPT-5.5 · Codex | SWE-Bench Pro: GPT-5.5/xhigh