

Dense-intelligence research


Continuous proposal—execution—verification over durable state

Invention

Long horizon

Verifiable

A Runtime architecture and fixed-model evolution



B Public evidence in task-native units

SOL-ExecBench
GPT-5.5 · Codex
Global #6 · 2×
#1 · 7 top-3
B200 · 101 kernels
Two head-to-head wins over Recursive

nanochat · B200
GPT-5.5 · Codex
0.9636 BPB
5 min · 1×B200 · 426 attempts
Human SOTA: 0.9646

nanochat · H100
GPT-5.5 · Codex
0.9855 BPB
5 min · 1×H100 · 37 mechanisms
Human SOTA: 0.9879

nanoGPT speedrun
GPT-5.5 · Codex
79.77 seconds
8×H100 · N=10
Same-device human #83: 80.18s

AARRI-Bench
GPT-5.5 · Codex
63/82 · 76.8%
82 research-intern tasks
Paper-reported best: 68.3%

Arbor comparison
GPT-5.5 · Codex
28.0 gap
Math-Reasoning Data
Arbor 20.83 · Claude Code 8.33 · Codex 6.25

41 de-duplicated research artifacts across 6 programs

35 manuscripts · 6 drafts

Architecture is schematic; result cards retain task-native units and are not cross-normalized.

All six runs: GPT-5.5 · Codex