


Dense-intelligence research

Continuous proposal—execution—verification over durable state

Invention Long horizon Verifiable



Benchmark results and full-suite software evaluation

SWE-Bench Pro · 731 tasks
GPT-5.5/xhigh

59% → 78%

Direct Copilot → Argus
1.41× aggregate tokens

SOL-ExecBench
Global #6 · 2× #1 · 7 top-3
B200 · 101 kernels
Two head-to-head wins over Recursive

nanochat · B200
0.9636 BPB
5 min · 1×B200 · 426 attempts
Human SOTA: 0.9646

nanochat · H100
0.9855 BPB
5 min · 1×H100 · 37 mechanisms
Human SOTA: 0.9879

nanoGPT speedrun
79.77 seconds
8×H100 · N=10
Same-device human #83: 80.18s

AARRI-Bench
63/82 · 76.8%
82 research-intern tasks
Paper-reported best: 68.3%

Math-Reasoning Data
28.0 gap
Math-Reasoning Data
Arbor 20.83 · Claude Code 8.33 · Codex 6.25

41 de-duplicated research papers across 6 programs

35 manuscripts · 6 working papers

Architecture is schematic; result cards retain task-native units and are not cross-normalized.

Public arenas: GPT-5.5 · Codex | SWE-Bench Pro: GPT-5.5/xhigh