

SWE-Bench Pro: accuracy gain and runtime convergence

One continuous 731-task evaluation · GPT-5.5/xhigh · Skill and Wiki state updated throughout the run

(a) Full-suite comparison

Accuracy on 731 tasks

Direct Copilot

≈59%

Argus

≈78%

+19 percentage points

Normalized aggregate tokens

Direct Copilot

1.00x

Argus

1.41x

Only the aggregate Token ratio is available;
Copilot per-wave Token/time was not logged.

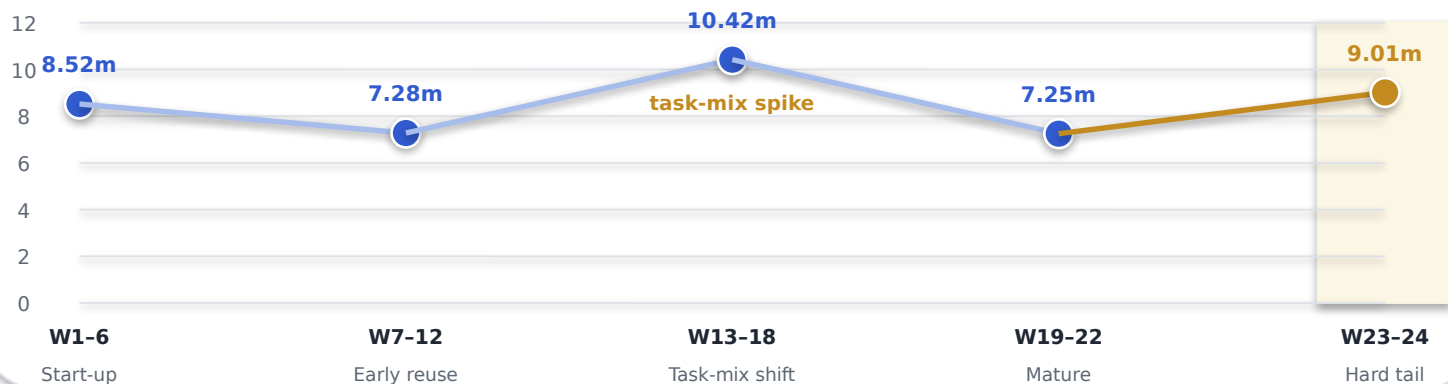
(b) Argus solve input per task

Start-up → mature: -21%



(c) Argus active workflow time per task

Start-up → mature: -15%



Task-weighted window means. W13 and W15 were deferred; W23-24 are retained as hard-tail stress. Queue, workspace/container, verifier, infrastructure-retry, and post-task maintenance time are excluded.